

Logistic Regression with Python and Scikit-Learn

In this project, I implement Logistic Regression with Python and Scikit-Learn. I build a classifier to predict whether or not it will rain tomorrow in Australia by training a binary classification model using Logistic Regression. I have used the **Rain in Australia** dataset downloaded from the Kaggle website for this project.

Table of Contents

The table of contents for this project is as follows:-

1. Introduction to Logistic Regression
2. Logistic Regression intuition
3. The problem statement
4. Dataset description
5. Import libraries
6. Import dataset
7. Exploratory data analysis
8. Declare feature vector and target variable
9. Split data into separate training and test set
10. Feature engineering
11. Feature scaling
12. Model training
13. Predict results
14. Check accuracy score
15. Confusion matrix
16. Classification metrics
17. Adjusting the threshold level
18. ROC - AUC
19. Recursive feature elimination
20. k-Fold Cross Validation
21. Hyperparameter optimization using GridSearch CV
22. Results and conclusion

1. Introduction to Logistic Regression

When data scientists may come across a new classification problem, the first algorithm that may come across their mind is **Logistic Regression**. It is a supervised learning classification algorithm which is used to predict observations to a discrete set of classes. Practically, it is used to classify observations into different categories. Hence, its output is discrete in nature. **Logistic Regression** is also called **Logit Regression**. It is one of the most simple, straightforward and versatile classification algorithms which is used to solve classification problems.

2. Logistic Regression intuition

In statistics, the **Logistic Regression model** is a widely used statistical model which is primarily used for classification purposes. It means that given a set of observations, Logistic Regression algorithm helps us to classify these observations into two or more discrete classes. So, the target variable is discrete in nature.

Logistic Regression algorithm works by implementing a linear equation with independent or explanatory variables to predict a response value. This predicted response value, denoted by z is then converted into a probability value that lie between 0 and 1. We use the **sigmoid function** in order to map predicted values to probability values. This sigmoid function then maps any real value into a probability value between 0 and 1.

The sigmoid function returns a probability value between 0 and 1. This probability value is then mapped to a discrete class which is either "0" or "1". In order to map this probability value to a discrete class (pass/fail, yes/no, true/false), we select a threshold value. This threshold value is called **Decision boundary**. Above this threshold value, we will map the probability values into class 1 and below which we will map values into class 0.

Mathematically, it can be expressed as follows:-

$$p \geq 0.5 \Rightarrow \text{class} = 1$$

$$p < 0.5 \Rightarrow \text{class} = 0$$

Generally, the decision boundary is set to 0.5. So, if the probability value is 0.8 (> 0.5), we will map this observation to class 1. Similarly, if the probability value is 0.2 (< 0.5), we will map this observation to class 0.

We can use our knowledge of **sigmoid function** and **decision boundary** to write a prediction function. A prediction function in logistic regression returns the probability of the observation being positive, **Yes** or **True**. We call this as **class 1** and it is denoted by $P(\text{class} = 1)$. If the probability inches closer to one, then we will be more confident about our model that the observation is in class 1.

Logistic regression intuition is discussed in depth in the readme document.

3. The problem statement

In this project, I try to answer the question that whether or not it will rain tomorrow in Australia. I implement Logistic Regression with Python and Scikit-Learn.

To answer the question, I build a classifier to predict whether or not it will rain tomorrow in Australia by training a binary classification model using Logistic Regression. I have used the **Rain in Australia** dataset downloaded from the Kaggle website for this project.

4. Dataset description

I have used the **Rain in Australia** data set downloaded from the Kaggle website.

I have downloaded this data set from the Kaggle website. The data set can be found at the following url:-

<https://www.kaggle.com/jsphyg/weather-dataset-rattle-package>

This dataset contains daily weather observations from numerous Australian weather stations.

5. Import libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings

warnings.filterwarnings('ignore')
```

6. Import dataset

```
data = 'C:/datasets/weatherAUS.csv'

df = pd.read_csv(data)
```

7. Exploratory data analysis

Now, I will explore the data to gain insights about the data.

```
# view dimensions of dataset

df.shape

(142193, 24)
```

We can see that there are 142193 instances and 24 variables in the data set.

```
# preview the dataset

df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation
Sunshine \						
0	2008-12-01	Albury	13.4	22.9	0.6	NaN
NaN						
1	2008-12-02	Albury	7.4	25.1	0.0	NaN
NaN						
2	2008-12-03	Albury	12.9	25.7	0.0	NaN
NaN						
3	2008-12-04	Albury	9.2	28.0	0.0	NaN
NaN						
4	2008-12-05	Albury	17.5	32.3	1.0	NaN
NaN						

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity3pm \
0	W	44.0	W	...	22.0
1	WNW	44.0	NNW	...	25.0
2	WSW	46.0	W	...	30.0
3	NE	24.0	SE	...	16.0
4	W	41.0	ENE	...	33.0

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm
RainToday \						
0	1007.7	1007.1	8.0	NaN	16.9	21.8
No						
1	1010.6	1007.8	NaN	NaN	17.2	24.3
No						
2	1007.6	1008.7	NaN	2.0	21.0	23.2
No						
3	1017.6	1012.8	NaN	NaN	18.1	26.5
No						
4	1010.8	1006.0	7.0	8.0	17.8	29.7
No						

	RISK_MM	RainTomorrow
0	0.0	No
1	0.0	No
2	0.0	No
3	1.0	No
4	0.2	No

```
[5 rows x 24 columns]
```

```
col_names = df.columns
```

```
col_names
```

```
Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall',
      'Evaporation',
      'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am',
      'WindDir3pm',
```

```

        'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
        'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm',
        'Temp9am',
        'Temp3pm', 'RainToday', 'RISK_MM', 'RainTomorrow'],
        dtype='object')

```

Drop RISK_MM variable

It is given in the dataset description, that we should drop the **RISK_MM** feature variable from the dataset description. So, we should drop it as follows-

```
df.drop(['RISK_MM'], axis=1, inplace=True)
```

```
# view summary of dataset
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 142193 entries, 0 to 142192
Data columns (total 23 columns):
Date                142193 non-null object
Location            142193 non-null object
MinTemp             141556 non-null float64
MaxTemp             141871 non-null float64
Rainfall            140787 non-null float64
Evaporation         81350 non-null float64
Sunshine            74377 non-null float64
WindGustDir         132863 non-null object
WindGustSpeed       132923 non-null float64
WindDir9am          132180 non-null object
WindDir3pm          138415 non-null object
WindSpeed9am        140845 non-null float64
WindSpeed3pm        139563 non-null float64
Humidity9am         140419 non-null float64
Humidity3pm         138583 non-null float64
Pressure9am         128179 non-null float64
Pressure3pm         128212 non-null float64
Cloud9am            88536 non-null float64
Cloud3pm            85099 non-null float64
Temp9am             141289 non-null float64
Temp3pm             139467 non-null float64
RainToday           140787 non-null object
RainTomorrow        142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.0+ MB

```

Types of variables

In this section, I segregate the dataset into categorical and numerical variables. There are a mixture of categorical and numerical variables in the dataset. Categorical variables have data type object. Numerical variables have data type float64.

First of all, I will find categorical variables.

```
# find categorical variables
```

```
categorical = [var for var in df.columns if df[var].dtype=='O']
```

```
print('There are {} categorical variables\n'.format(len(categorical)))
```

```
print('The categorical variables are :', categorical)
```

There are 7 categorical variables

The categorical variables are : ['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']

```
# view the categorical variables
```

```
df[categorical].head()
```

	Date	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday	\
0	2008-12-01	Albury	W	W	WNW	No	
1	2008-12-02	Albury	WNW	NNW	WSW	No	
2	2008-12-03	Albury	WSW	W	WSW	No	
3	2008-12-04	Albury	NE	SE	E	No	
4	2008-12-05	Albury	W	ENE	NW	No	

	RainTomorrow
0	No
1	No
2	No
3	No
4	No

Summary of categorical variables

- There is a date variable. It is denoted by `Date` column.
- There are 6 categorical variables. These are given by `Location`, `WindGustDir`, `WindDir9am`, `WindDir3pm`, `RainToday` and `RainTomorrow`.
- There are two binary categorical variables - `RainToday` and `RainTomorrow`.
- `RainTomorrow` is the target variable.

Explore problems within categorical variables

First, I will explore the categorical variables.

Missing values in categorical variables

```
# check missing values in categorical variables

df[categorical].isnull().sum()

Date          0
Location      0
WindGustDir    9330
WindDir9am    10013
WindDir3pm    3778
RainToday     1406
RainTomorrow  0
dtype: int64

# print categorical variables containing missing values

cat1 = [var for var in categorical if df[var].isnull().sum()!=0]

print(df[cat1].isnull().sum())

WindGustDir    9330
WindDir9am     10013
WindDir3pm     3778
RainToday      1406
dtype: int64
```

We can see that there are only 4 categorical variables in the dataset which contains missing values. These are `WindGustDir`, `WindDir9am`, `WindDir3pm` and `RainToday`.

Frequency counts of categorical variables

Now, I will check the frequency counts of categorical variables.

```
# view frequency of categorical variables

for var in categorical:

    print(df[var].value_counts())

2014-10-12    49
2017-01-15    49
2013-10-02    49
2014-07-15    49
2014-02-19    49
2016-08-21    49
2014-07-03    49
```

2016-10-21	49
2013-03-11	49
2017-02-08	49
2014-11-17	49
2013-04-25	49
2014-11-19	49
2014-08-30	49
2014-01-07	49
2013-04-10	49
2017-03-16	49
2013-09-04	49
2016-08-16	49
2016-10-19	49
2014-08-20	49
2017-05-12	49
2014-01-16	49
2016-07-22	49
2017-01-22	49
2013-09-25	49
2013-06-02	49
2016-07-06	49
2014-04-21	49
2013-10-16	49
	..
2007-11-23	1
2008-01-15	1
2007-12-22	1
2007-11-08	1
2007-11-29	1
2008-01-29	1
2008-01-06	1
2007-11-02	1
2007-12-25	1
2008-01-28	1
2007-12-08	1
2007-11-09	1
2008-01-05	1
2007-11-26	1
2007-11-10	1
2007-11-20	1
2008-01-14	1
2007-12-03	1
2008-01-12	1
2007-11-03	1
2007-12-02	1
2008-01-31	1
2007-12-01	1
2007-11-06	1
2007-11-27	1


```
2007-12-19      1
2007-11-19      1
2007-12-30      1
2007-12-23      1
2008-01-09      1
Name: Date, Length: 3436, dtype: int64
Canberra        3418
Sydney          3337
Perth           3193
Darwin          3192
Hobart          3188
Brisbane        3161
Adelaide        3090
Bendigo         3034
Townsville      3033
AliceSprings    3031
MountGambier    3030
Launceston      3028
Ballarat        3028
Albany          3016
Albury          3011
PerthAirport    3009
MelbourneAirport 3009
Mildura         3007
SydneyAirport   3005
Nuriootpa       3002
Sale            3000
Watsonia        2999
Tuggeranong     2998
Portland        2996
Woomera         2990
Cairns          2988
Cobar           2988
Wollongong      2983
GoldCoast       2980
WaggaWagga      2976
NorfolkIsland   2964
Penrith         2964
Newcastle       2955
SalmonGums      2955
CoffsHarbour    2953
Witchcliffe     2952
Richmond        2951
Dartmoor        2943
NorahHead       2929
BadgerysCreek   2928
MountGinini     2907
Moree           2854
Walpole         2819
```

PearceRAAF	2762
Williamstown	2553
Melbourne	2435
Nhil	1569
Katherine	1559
Uluru	1521

Name: Location, dtype: int64

W	9780
SE	9309
E	9071
N	9033
SSE	8993
S	8949
WSW	8901
SW	8797
SSW	8610
WNW	8066
NW	8003
ENE	7992
ESE	7305
NE	7060
NNW	6561
NNE	6433

Name: WindGustDir, dtype: int64

N	11393
SE	9162
E	9024
SSE	8966
NW	8552
S	8493
W	8260
SW	8237
NNE	7948
NNW	7840
ENE	7735
ESE	7558
NE	7527
SSW	7448
WNW	7194
WSW	6843

Name: WindDir9am, dtype: int64

SE	10663
W	9911
S	9598
WSW	9329
SW	9182
SSE	9142
N	8667
WNW	8656

```
NW      8468
ESE      8382
E        8342
NE       8164
SSW      8010
NNW      7733
ENE      7724
NNE      6444
Name: WindDir3pm, dtype: int64
No       109332
Yes      31455
Name: RainToday, dtype: int64
No       110316
Yes      31877
Name: RainTomorrow, dtype: int64
```

```
# view frequency distribution of categorical variables
```

```
for var in categorical:
```

```
    print(df[var].value_counts()/np.float(len(df)))
```

```
2014-10-12    0.000345
2017-01-15    0.000345
2013-10-02    0.000345
2014-07-15    0.000345
2014-02-19    0.000345
2016-08-21    0.000345
2014-07-03    0.000345
2016-10-21    0.000345
2013-03-11    0.000345
2017-02-08    0.000345
2014-11-17    0.000345
2013-04-25    0.000345
2014-11-19    0.000345
2014-08-30    0.000345
2014-01-07    0.000345
2013-04-10    0.000345
2017-03-16    0.000345
2013-09-04    0.000345
2016-08-16    0.000345
2016-10-19    0.000345
2014-08-20    0.000345
2017-05-12    0.000345
2014-01-16    0.000345
2016-07-22    0.000345
2017-01-22    0.000345
2013-09-25    0.000345
2013-06-02    0.000345
2016-07-06    0.000345
```

2014-04-21	0.000345
2013-10-16	0.000345
	...
2007-11-23	0.000007
2008-01-15	0.000007
2007-12-22	0.000007
2007-11-08	0.000007
2007-11-29	0.000007
2008-01-29	0.000007
2008-01-06	0.000007
2007-11-02	0.000007
2007-12-25	0.000007
2008-01-28	0.000007
2007-12-08	0.000007
2007-11-09	0.000007
2008-01-05	0.000007
2007-11-26	0.000007
2007-11-10	0.000007
2007-11-20	0.000007
2008-01-14	0.000007
2007-12-03	0.000007
2008-01-12	0.000007
2007-11-03	0.000007
2007-12-02	0.000007
2008-01-31	0.000007
2007-12-01	0.000007
2007-11-06	0.000007
2007-11-27	0.000007
2007-12-19	0.000007
2007-11-19	0.000007
2007-12-30	0.000007
2007-12-23	0.000007
2008-01-09	0.000007
Name: Date, Length: 3436, dtype: float64	
Canberra	0.024038
Sydney	0.023468
Perth	0.022455
Darwin	0.022448
Hobart	0.022420
Brisbane	0.022230
Adelaide	0.021731
Bendigo	0.021337
Townsville	0.021330
AliceSprings	0.021316
MountGambier	0.021309
Launceston	0.021295
Ballarat	0.021295
Albany	0.021211
Albury	0.021175

PerthAirport	0.021161
MelbourneAirport	0.021161
Mildura	0.021147
SydneyAirport	0.021133
Nuriootpa	0.021112
Sale	0.021098
Watsonia	0.021091
Tuggeranong	0.021084
Portland	0.021070
Woomera	0.021028
Cairns	0.021014
Cobar	0.021014
Wollongong	0.020979
GoldCoast	0.020957
WaggaWagga	0.020929
NorfolkIsland	0.020845
Penrith	0.020845
Newcastle	0.020782
SalmonGums	0.020782
CoffsHarbour	0.020768
Witchcliffe	0.020761
Richmond	0.020753
Dartmoor	0.020697
NorahHead	0.020599
BadgerysCreek	0.020592
MountGinini	0.020444
Moree	0.020071
Walpole	0.019825
PearceRAAF	0.019424
Williamstown	0.017954
Melbourne	0.017125
Nhil	0.011034
Katherine	0.010964
Uluru	0.010697

Name: Location, dtype: float64

W	0.068780
SE	0.065467
E	0.063794
N	0.063526
SSE	0.063245
S	0.062936
WSW	0.062598
SW	0.061867
SSW	0.060552
WNW	0.056726
NW	0.056283
ENE	0.056205
ESE	0.051374
NE	0.049651

```

NNW      0.046142
NNE      0.045241
Name: WindGustDir, dtype: float64
N        0.080123
SE       0.064434
E        0.063463
SSE      0.063055
NW       0.060144
S        0.059729
W        0.058090
SW       0.057928
NNE      0.055896
NNW      0.055136
ENE      0.054398
ESE      0.053153
NE       0.052935
SSW      0.052380
WNW      0.050593
WSW      0.048125
Name: WindDir9am, dtype: float64
SE       0.074990
W        0.069701
S        0.067500
WSW      0.065608
SW       0.064574
SSE      0.064293
N        0.060952
WNW      0.060875
NW       0.059553
ESE      0.058948
E        0.058667
NE       0.057415
SSW      0.056332
NNW      0.054384
ENE      0.054321
NNE      0.045319
Name: WindDir3pm, dtype: float64
No       0.768899
Yes      0.221213
Name: RainToday, dtype: float64
No       0.775819
Yes      0.224181
Name: RainTomorrow, dtype: float64

```

Number of labels: cardinality

The number of labels within a categorical variable is known as **cardinality**. A high number of labels within a variable is known as **high cardinality**. High cardinality may pose some serious problems in the machine learning model. So, I will check for high cardinality.

```
# check for cardinality in categorical variables

for var in categorical:

    print(var, ' contains ', len(df[var].unique()), ' labels')

Date contains 3436 labels
Location contains 49 labels
WindGustDir contains 17 labels
WindDir9am contains 17 labels
WindDir3pm contains 17 labels
RainToday contains 3 labels
RainTomorrow contains 2 labels
```

We can see that there is a `Date` variable which needs to be preprocessed. I will do preprocessing in the following section.

All the other variables contain relatively smaller number of variables.

Feature Engineering of Date Variable

```
df['Date'].dtypes
dtype('O')
```

We can see that the data type of `Date` variable is object. I will parse the date currently coded as object into datetime format.

```
# parse the dates, currently coded as strings, into datetime format

df['Date'] = pd.to_datetime(df['Date'])

# extract year from date

df['Year'] = df['Date'].dt.year

df['Year'].head()

0    2008
1    2008
2    2008
3    2008
4    2008
Name: Year, dtype: int64

# extract month from date

df['Month'] = df['Date'].dt.month

df['Month'].head()
```

```
0    12
1    12
2    12
3    12
4    12
```

```
Name: Month, dtype: int64
```

```
# extract day from date
```

```
df['Day'] = df['Date'].dt.day
```

```
df['Day'].head()
```

```
0    1
1    2
2    3
3    4
4    5
```

```
Name: Day, dtype: int64
```

```
# again view the summary of dataset
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 142193 entries, 0 to 142192
```

```
Data columns (total 26 columns):
```

Date	142193	non-null	datetime64[ns]
Location	142193	non-null	object
MinTemp	141556	non-null	float64
MaxTemp	141871	non-null	float64
Rainfall	140787	non-null	float64
Evaporation	81350	non-null	float64
Sunshine	74377	non-null	float64
WindGustDir	132863	non-null	object
WindGustSpeed	132923	non-null	float64
WindDir9am	132180	non-null	object
WindDir3pm	138415	non-null	object
WindSpeed9am	140845	non-null	float64
WindSpeed3pm	139563	non-null	float64
Humidity9am	140419	non-null	float64
Humidity3pm	138583	non-null	float64
Pressure9am	128179	non-null	float64
Pressure3pm	128212	non-null	float64
Cloud9am	88536	non-null	float64
Cloud3pm	85099	non-null	float64
Temp9am	141289	non-null	float64
Temp3pm	139467	non-null	float64
RainToday	140787	non-null	object
RainTomorrow	142193	non-null	object


```

Year          142193 non-null int64
Month         142193 non-null int64
Day           142193 non-null int64
dtypes: datetime64[ns](1), float64(16), int64(3), object(6)
memory usage: 28.2+ MB

```

We can see that there are three additional columns created from `Date` variable. Now, I will drop the original `Date` variable from the dataset.

```
# drop the original Date variable
```

```
df.drop('Date', axis=1, inplace = True)
```

```
# preview the dataset again
```

```
df.head()
```

	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine
0	Albury	13.4	22.9	0.6	NaN	NaN
1	Albury	7.4	25.1	0.0	NaN	NaN
2	Albury	12.9	25.7	0.0	NaN	NaN
3	Albury	9.2	28.0	0.0	NaN	NaN
4	Albury	17.5	32.3	1.0	NaN	NaN

	WindGustDir	WindGustSpeed	WindDir9am	WindDir3pm	...	Pressure3pm	Cloud9am
0	W	44.0	W	WNW	...	1007.1	8.0
1	WNW	44.0	NNW	WSW	...	1007.8	NaN
2	WSW	46.0	W	WSW	...	1008.7	NaN
3	NE	24.0	SE	E	...	1012.8	NaN
4	W	41.0	ENE	NW	...	1006.0	7.0

	Temp9am	Temp3pm	RainToday	RainTomorrow	Year	Month	Day
0	16.9	21.8	No	No	2008	12	1
1	17.2	24.3	No	No	2008	12	2
2	21.0	23.2	No	No	2008	12	3
3	18.1	26.5	No	No	2008	12	4
4	17.8	29.7	No	No	2008	12	5

```
[5 rows x 25 columns]
```

Now, we can see that the `Date` variable has been removed from the dataset.

Explore Categorical Variables

Now, I will explore the categorical variables one by one.

```
# find categorical variables
categorical = [var for var in df.columns if df[var].dtype=='O']
print('There are {} categorical variables\n'.format(len(categorical)))
print('The categorical variables are :', categorical)

There are 6 categorical variables

The categorical variables are : ['Location', 'WindGustDir',
'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']
```

We can see that there are 6 categorical variables in the dataset. The `Date` variable has been removed. First, I will check missing values in categorical variables.

```
# check for missing values in categorical variables
df[categorical].isnull().sum()

Location          0
WindGustDir       9330
WindDir9am        10013
WindDir3pm        3778
RainToday         1406
RainTomorrow      0
dtype: int64
```

We can see that `WindGustDir`, `WindDir9am`, `WindDir3pm`, `RainToday` variables contain missing values. I will explore these variables one by one.

Explore Location variable

```
# print number of labels in Location variable
print('Location contains', len(df.Location.unique()), 'labels')

Location contains 49 labels
```

```
# check labels in location variable
```

```
df.Location.unique()
```

```
array(['Albury', 'BadgerysCreek', 'Cobar', 'CoffsHarbour', 'Moree',  
      'Newcastle', 'NorahHead', 'NorfolkIsland', 'Penrith',  
      'Richmond',  
      'Sydney', 'SydneyAirport', 'WaggaWagga', 'Williamtown',  
      'Wollongong', 'Canberra', 'Tuggeranong', 'MountGinini',  
      'Ballarat',  
      'Bendigo', 'Sale', 'MelbourneAirport', 'Melbourne', 'Mildura',  
      'Nhil', 'Portland', 'Watsonia', 'Dartmoor', 'Brisbane',  
      'Cairns',  
      'GoldCoast', 'Townsville', 'Adelaide', 'MountGambier',  
      'Nuriootpa',  
      'Woomera', 'Albany', 'Witchcliffe', 'PearceRAAF',  
      'PerthAirport',  
      'Perth', 'SalmonGums', 'Walpole', 'Hobart', 'Launceston',  
      'AliceSprings', 'Darwin', 'Katherine', 'Uluru'], dtype=object)
```

```
# check frequency distribution of values in Location variable
```

```
df.Location.value_counts()
```

Canberra	3418
Sydney	3337
Perth	3193
Darwin	3192
Hobart	3188
Brisbane	3161
Adelaide	3090
Bendigo	3034
Townsville	3033
AliceSprings	3031
MountGambier	3030
Launceston	3028
Ballarat	3028
Albany	3016
Albury	3011
PerthAirport	3009
MelbourneAirport	3009
Mildura	3007
SydneyAirport	3005
Nuriootpa	3002
Sale	3000
Watsonia	2999
Tuggeranong	2998
Portland	2996
Woomera	2990
Cairns	2988

Cobar	2988
Wollongong	2983
GoldCoast	2980
WaggaWagga	2976
NorfolkIsland	2964
Penrith	2964
Newcastle	2955
SalmonGums	2955
CoffsHarbour	2953
Witchcliffe	2952
Richmond	2951
Dartmoor	2943
NorahHead	2929
BadgerysCreek	2928
MountGinini	2907
Moree	2854
Walpole	2819
PearceRAAF	2762
Williamstown	2553
Melbourne	2435
Nhil	1569
Katherine	1559
Uluru	1521

Name: Location, dtype: int64

```
# let's do One Hot Encoding of Location variable
# get k-1 dummy variables after One Hot Encoding
# preview the dataset with head() method
```

```
pd.get_dummies(df.Location, drop_first=True).head()
```

	Albany	Albury	AliceSprings	BadgerysCreek	Ballarat	Bendigo
0	0	1	0	0	0	0
1	0	1	0	0	0	0
2	0	1	0	0	0	0
3	0	1	0	0	0	0
4	0	1	0	0	0	0

	Cairns	Canberra	Cobar	...	Townsville	Tuggeranong	Uluru	\
0	0	0	0	...	0	0	0	
1	0	0	0	...	0	0	0	
2	0	0	0	...	0	0	0	
3	0	0	0	...	0	0	0	
4	0	0	0	...	0	0	0	

	WaggaWagga	Walpole	Watsonia	Williamstown	Witchcliffe	Wollongong
\						
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

	Woomera
0	0
1	0
2	0
3	0
4	0

[5 rows x 48 columns]

Explore WindGustDir variable

```
# print number of labels in WindGustDir variable
```

```
print('WindGustDir contains', len(df['WindGustDir'].unique()),  
      'labels')
```

WindGustDir contains 17 labels

```
# check labels in WindGustDir variable
```

```
df['WindGustDir'].unique()
```

```
array(['W', 'WNW', 'WSW', 'NE', 'NNW', 'N', 'NNE', 'SW', 'ENE', 'SSE',  
      'S', 'NW', 'SE', 'ESE', nan, 'E', 'SSW'], dtype=object)
```

```
# check frequency distribution of values in WindGustDir variable
```

```
df.WindGustDir.value_counts()
```

```
W      9780  
SE     9309  
E      9071  
N      9033  
SSE    8993  
S      8949  
WSW    8901  
SW     8797
```

```
SSW      8610
WNW      8066
NW        8003
ENE       7992
ESE       7305
NE        7060
NNW       6561
NNE       6433
```

```
Name: WindGustDir, dtype: int64
```

```
# let's do One Hot Encoding of WindGustDir variable
# get k-1 dummy variables after One Hot Encoding
# also add an additional dummy variable to indicate there was missing data
# preview the dataset with head() method
```

```
pd.get_dummies(df.WindGustDir, drop_first=True, dummy_na=True).head()
```

	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW
NaN															
0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
0															
1	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0
0															
2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
0															
3	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0															
4	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
0															

```
# sum the number of 1s per boolean variable over the rows of the dataset
# it will tell us how many observations we have for each category
```

```
pd.get_dummies(df.WindGustDir, drop_first=True, dummy_na=True).sum(axis=0)
```

```
ENE      7992
ESE      7305
N        9033
NE       7060
NNE      6433
NNW      6561
NW       8003
S        8949
SE       9309
SSE      8993
SSW      8610
SW       8797
```

```
W      9780
WNW    8066
WSW    8901
NaN    9330
dtype: int64
```

We can see that there are 9330 missing values in WindGustDir variable.

Explore WindDir9am variable

```
# print number of labels in WindDir9am variable
print('WindDir9am contains', len(df['WindDir9am'].unique()), 'labels')
WindDir9am contains 17 labels

# check labels in WindDir9am variable
df['WindDir9am'].unique()
array(['W', 'NNW', 'SE', 'ENE', 'SW', 'SSE', 'S', 'NE', nan, 'SSW',
       'N', 'WSW', 'ESE', 'E', 'NW', 'WNW', 'NNE'], dtype=object)

# check frequency distribution of values in WindDir9am variable
df['WindDir9am'].value_counts()
N      11393
SE      9162
E       9024
SSE     8966
NW      8552
S       8493
W       8260
SW      8237
NNE     7948
NNW     7840
ENE     7735
ESE     7558
NE      7527
SSW     7448
WNW     7194
WSW     6843
Name: WindDir9am, dtype: int64

# let's do One Hot Encoding of WindDir9am variable
# get k-1 dummy variables after One Hot Encoding
# also add an additional dummy variable to indicate there was missing
data
# preview the dataset with head() method
```

```
pd.get_dummies(df.WindDir9am, drop_first=True, dummy_na=True).head()
```

	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW
NaN															
0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
1	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
3	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```
# sum the number of 1s per boolean variable over the rows of the dataset
# it will tell us how many observations we have for each category
```

```
pd.get_dummies(df.WindDir9am, drop_first=True,
dummy_na=True).sum(axis=0)
```

```
ENE      7735
ESE      7558
N        11393
NE       7527
NNE      7948
NNW      7840
NW       8552
S        8493
SE       9162
SSE      8966
SSW      7448
SW       8237
W        8260
WNW      7194
WSW      6843
NaN      10013
dtype: int64
```

We can see that there are 10013 missing values in the `WindDir9am` variable.

Explore WindDir3pm variable

```
# print number of labels in WindDir3pm variable
```

```
print('WindDir3pm contains', len(df['WindDir3pm'].unique()), 'labels')
```

```
WindDir3pm contains 17 labels
```


[illegible]

```
# sum the number of 1s per boolean variable over the rows of the dataset  
# it will tell us how many observations we have for each category
```

```
pd.get_dummies(df.WindDir3pm, drop_first=True,  
dummy_na=True).sum(axis=0)
```

```
ENE      7724  
ESE      8382  
N        8667  
NE       8164  
NNE      6444  
NNW      7733  
NW       8468  
S        9598  
SE       10663  
SSE      9142  
SSW      8010  
SW       9182  
W        9911  
WNW      8656  
WSW      9329  
NaN      3778  
dtype: int64
```

There are 3778 missing values in the `WindDir3pm` variable.

Explore RainToday variable

```
# print number of labels in RainToday variable
```

```
print('RainToday contains', len(df['RainToday'].unique()), 'labels')
```

```
RainToday contains 3 labels
```

```
# check labels in WindGustDir variable
```

```
df['RainToday'].unique()
```

```
array(['No', 'Yes', nan], dtype=object)
```

```
# check frequency distribution of values in WindGustDir variable
```

```
df.RainToday.value_counts()
```

```
No      109332
```

```
Yes      31455
```

```
Name: RainToday, dtype: int64
```

```
# let's do One Hot Encoding of RainToday variable
```

```
# get k-1 dummy variables after One Hot Encoding
```

```
# also add an additional dummy variable to indicate there was missing data
# preview the dataset with head() method
```

```
pd.get_dummies(df.RainToday, drop_first=True, dummy_na=True).head()
```

	Yes	NaN
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

```
# sum the number of 1s per boolean variable over the rows of the dataset
# it will tell us how many observations we have for each category
```

```
pd.get_dummies(df.RainToday, drop_first=True, dummy_na=True).sum(axis=0)
```

```
Yes      31455
NaN       1406
dtype: int64
```

There are 1406 missing values in the `RainToday` variable.

Explore Numerical Variables

```
# find numerical variables
```

```
numerical = [var for var in df.columns if df[var].dtype!='O']
print('There are {} numerical variables\n'.format(len(numerical)))
print('The numerical variables are :', numerical)
```

There are 19 numerical variables

```
The numerical variables are : ['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm', 'Year', 'Month', 'Day']
```

```
# view the numerical variables
```

```
df[numerical].head()
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	13.4	22.9	0.6	NaN	NaN	44.0	
1	7.4	25.1	0.0	NaN	NaN	44.0	

2	12.9	25.7	0.0	NaN	NaN	46.0
3	9.2	28.0	0.0	NaN	NaN	24.0
4	17.5	32.3	1.0	NaN	NaN	41.0
WindSpeed9am WindSpeed3pm Humidity9am Humidity3pm						
Pressure9am \						
0	20.0	24.0	71.0	22.0	1007.7	
1	4.0	22.0	44.0	25.0	1010.6	
2	19.0	26.0	38.0	30.0	1007.6	
3	11.0	9.0	45.0	16.0	1017.6	
4	7.0	20.0	82.0	33.0	1010.8	
Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm Year Month Day						
0	1007.1	8.0	NaN	16.9	21.8	2008 12 1
1	1007.8	NaN	NaN	17.2	24.3	2008 12 2
2	1008.7	NaN	2.0	21.0	23.2	2008 12 3
3	1012.8	NaN	NaN	18.1	26.5	2008 12 4
4	1006.0	7.0	8.0	17.8	29.7	2008 12 5

Summary of numerical variables

- There are 16 numerical variables.
- These are given by MinTemp, MaxTemp, Rainfall, Evaporation, Sunshine, WindGustSpeed, WindSpeed9am, WindSpeed3pm, Humidity9am, Humidity3pm, Pressure9am, Pressure3pm, Cloud9am, Cloud3pm, Temp9am and Temp3pm.
- All of the numerical variables are of continuous type.

Explore problems within numerical variables

Now, I will explore the numerical variables.

Missing values in numerical variables

```
# check missing values in numerical variables
```

```
df[numerical].isnull().sum()
```

```
MinTemp      637
MaxTemp      322
```

```

Rainfall      1406
Evaporation   60843
Sunshine      67816
WindGustSpeed  9270
WindSpeed9am  1348
WindSpeed3pm  2630
Humidity9am   1774
Humidity3pm   3610
Pressure9am   14014
Pressure3pm   13981
Cloud9am      53657
Cloud3pm      57094
Temp9am       904
Temp3pm       2726
Year          0
Month         0
Day           0
dtype: int64

```

We can see that all the 16 numerical variables contain missing values.

Outliers in numerical variables

```
# view summary statistics in numerical variables
```

```
print(round(df[numerical].describe()),2)
```

```

      MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine
WindGustSpeed \
count  141556.0  141871.0  140787.0      81350.0   74377.0
132923.0
mean      12.0     23.0       2.0         5.0     8.0
40.0
std       6.0      7.0       8.0         4.0     4.0
14.0
min      -8.0     -5.0       0.0         0.0     0.0
6.0
25%       8.0     18.0       0.0         3.0     5.0
31.0
50%      12.0     23.0       0.0         5.0     8.0
39.0
75%      17.0     28.0       1.0         7.0    11.0
48.0
max      34.0     48.0     371.0      145.0    14.0
135.0

      WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm
Pressure9am \
count    140845.0    139563.0    140419.0    138583.0
128179.0

```

mean	14.0	19.0	69.0	51.0
1018.0				
std	9.0	9.0	19.0	21.0
7.0				
min	0.0	0.0	0.0	0.0
980.0				
25%	7.0	13.0	57.0	37.0
1013.0				
50%	13.0	19.0	70.0	52.0
1018.0				
75%	19.0	24.0	83.0	66.0
1022.0				
max	130.0	87.0	100.0	100.0
1041.0				

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	
Year \						
count	128212.0	88536.0	85099.0	141289.0	139467.0	142193.0
mean	1015.0	4.0	5.0	17.0	22.0	2013.0
std	7.0	3.0	3.0	6.0	7.0	3.0
min	977.0	0.0	0.0	-7.0	-5.0	2007.0
25%	1010.0	1.0	2.0	12.0	17.0	2011.0
50%	1015.0	5.0	5.0	17.0	21.0	2013.0
75%	1020.0	7.0	7.0	22.0	26.0	2015.0
max	1040.0	9.0	9.0	40.0	47.0	2017.0

	Month	Day
count	142193.0	142193.0
mean	6.0	16.0
std	3.0	9.0
min	1.0	1.0
25%	3.0	8.0
50%	6.0	16.0
75%	9.0	23.0
max	12.0	31.0

On closer inspection, we can see that the Rainfall, Evaporation, WindSpeed9am and WindSpeed3pm columns may contain outliers.

I will draw boxplots to visualise outliers in the above variables.

```
# draw boxplots to visualize outliers

plt.figure(figsize=(15,10))

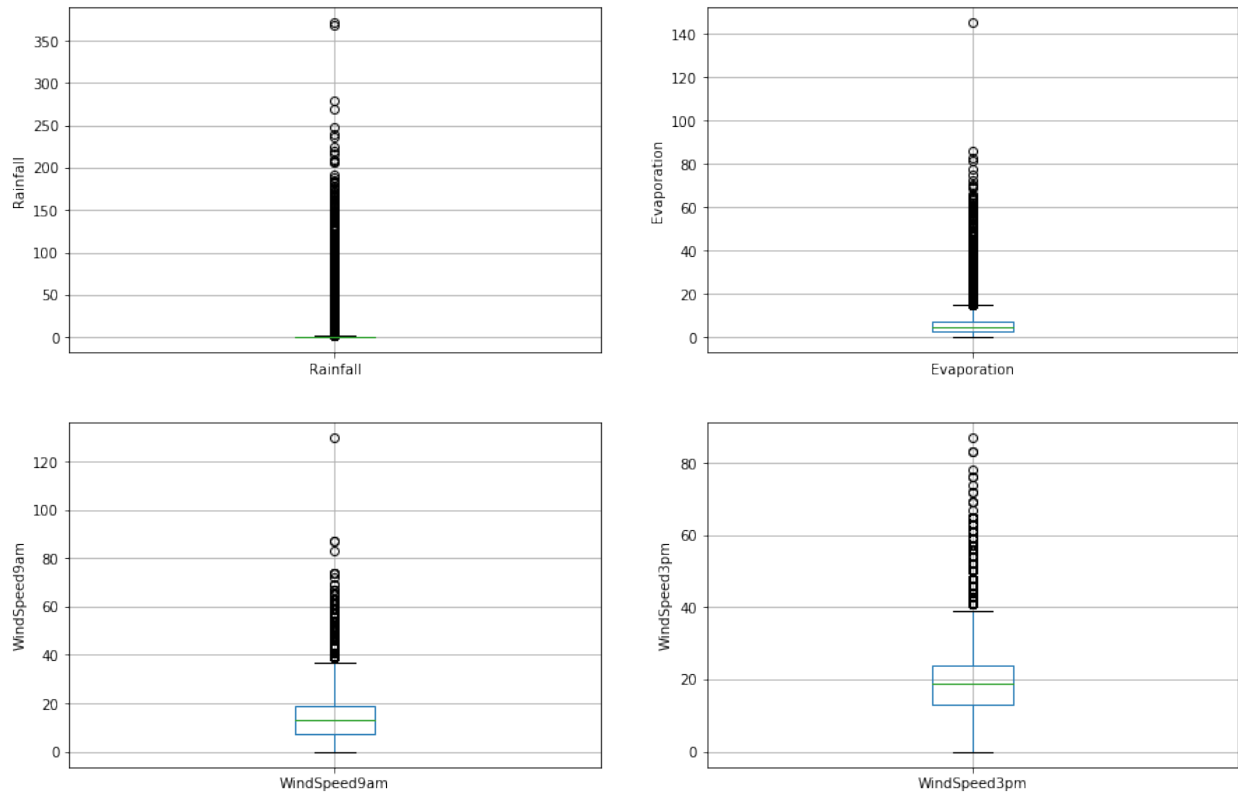
plt.subplot(2, 2, 1)
fig = df.boxplot(column='Rainfall')
fig.set_title('')
fig.set_ylabel('Rainfall')

plt.subplot(2, 2, 2)
fig = df.boxplot(column='Evaporation')
fig.set_title('')
fig.set_ylabel('Evaporation')

plt.subplot(2, 2, 3)
fig = df.boxplot(column='WindSpeed9am')
fig.set_title('')
fig.set_ylabel('WindSpeed9am')

plt.subplot(2, 2, 4)
fig = df.boxplot(column='WindSpeed3pm')
fig.set_title('')
fig.set_ylabel('WindSpeed3pm')

Text(0,0.5, 'WindSpeed3pm')
```



The above boxplots confirm that there are lot of outliers in these variables.

Check the distribution of variables

Now, I will plot the histograms to check distributions to find out if they are normal or skewed. If the variable follows normal distribution, then I will do **Extreme Value Analysis** otherwise if they are skewed, I will find IQR (Interquartile range).

```
# plot histogram to check distribution
```

```
plt.figure(figsize=(15,10))
```

```
plt.subplot(2, 2, 1)
fig = df.Rainfall.hist(bins=10)
fig.set_xlabel('Rainfall')
fig.set_ylabel('RainTomorrow')
```

```
plt.subplot(2, 2, 2)
fig = df.Evaporation.hist(bins=10)
fig.set_xlabel('Evaporation')
fig.set_ylabel('RainTomorrow')
```





```
plt.subplot(2, 2, 3)
```



```
fig = df.WindSpeed9am.hist(bins=10)
fig.set_xlabel('WindSpeed9am')
fig.set_ylabel('RainTomorrow')
```

```
plt.subplot(2, 2, 4)
fig = df.WindSpeed3pm.hist(bins=10)
fig.set_xlabel('WindSpeed3pm')
fig.set_ylabel('RainTomorrow')
```

```
Text(0,0.5,'RainTomorrow')
```



We can see that all the four variables are skewed. So, I will use interquartile range to find outliers.

```
# find outliers for Rainfall variable
```

```
IQR = df.Rainfall.quantile(0.75) - df.Rainfall.quantile(0.25)
Lower_fence = df.Rainfall.quantile(0.25) - (IQR * 3)
Upper_fence = df.Rainfall.quantile(0.75) + (IQR * 3)
print('Rainfall outliers are values < {lowerboundary} or >
{upperboundary}'.format(lowerboundary=Lower_fence,
upperboundary=Upper_fence))
```

```
Rainfall outliers are values < -2.4000000000000004 or > 3.2
```

For **Rainfall**, the minimum and maximum values are 0.0 and 371.0. So, the outliers are values > 3.2.

```
# find outliers for Evaporation variable
```

```
IQR = df.Evaporation.quantile(0.75) - df.Evaporation.quantile(0.25)
Lower_fence = df.Evaporation.quantile(0.25) - (IQR * 3)
Upper_fence = df.Evaporation.quantile(0.75) + (IQR * 3)
print('Evaporation outliers are values < {lowerboundary} or >
{upperboundary}'.format(lowerboundary=Lower_fence,
upperboundary=Upper_fence))
```

```
Evaporation outliers are values < -11.800000000000002 or >
21.800000000000004
```

For **Evaporation**, the minimum and maximum values are 0.0 and 145.0. So, the outliers are values > 21.8.

```
# find outliers for WindSpeed9am variable
```

```
IQR = df.WindSpeed9am.quantile(0.75) - df.WindSpeed9am.quantile(0.25)
Lower_fence = df.WindSpeed9am.quantile(0.25) - (IQR * 3)
Upper_fence = df.WindSpeed9am.quantile(0.75) + (IQR * 3)
print('WindSpeed9am outliers are values < {lowerboundary} or >
{upperboundary}'.format(lowerboundary=Lower_fence,
upperboundary=Upper_fence))
```

```
WindSpeed9am outliers are values < -29.0 or > 55.0
```

For **WindSpeed9am**, the minimum and maximum values are 0.0 and 130.0. So, the outliers are values > 55.0.

```
# find outliers for WindSpeed3pm variable
```

```
IQR = df.WindSpeed3pm.quantile(0.75) - df.WindSpeed3pm.quantile(0.25)
Lower_fence = df.WindSpeed3pm.quantile(0.25) - (IQR * 3)
Upper_fence = df.WindSpeed3pm.quantile(0.75) + (IQR * 3)
print('WindSpeed3pm outliers are values < {lowerboundary} or >
{upperboundary}'.format(lowerboundary=Lower_fence,
upperboundary=Upper_fence))
```

```
WindSpeed3pm outliers are values < -20.0 or > 57.0
```

For **WindSpeed3pm**, the minimum and maximum values are 0.0 and 87.0. So, the outliers are values > 57.0.

8. Declare feature vector and target variable

```
X = df.drop(['RainTomorrow'], axis=1)
y = df['RainTomorrow']
```

9. Split data into separate training and test set

```
# split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state = 0)

# check the shape of X_train and X_test

X_train.shape, X_test.shape

((113754, 24), (28439, 24))
```

10. Feature Engineering

Feature Engineering is the process of transforming raw data into useful features that help us to understand our model better and increase its predictive power. I will carry out feature engineering on different types of variables.

First, I will display the categorical and numerical variables again separately.

```
# check data types in X_train

X_train.dtypes

Location          object
MinTemp           float64
MaxTemp           float64
Rainfall          float64
Evaporation       float64
Sunshine          float64
WindGustDir       object
WindGustSpeed     float64
WindDir9am        object
WindDir3pm        object
WindSpeed9am      float64
WindSpeed3pm      float64
Humidity9am       float64
Humidity3pm       float64
Pressure9am       float64
Pressure3pm       float64
Cloud9am          float64
```

```

Cloud3pm          float64
Temp9am           float64
Temp3pm           float64
RainToday         object
Year              int64
Month             int64
Day               int64
dtype: object

# display categorical variables

categorical = [col for col in X_train.columns if X_train[col].dtypes
== 'O']

categorical

['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday']

# display numerical variables

numerical = [col for col in X_train.columns if X_train[col].dtypes !=
'O']

numerical

['MinTemp',
'MaxTemp',
'Rainfall',
'Evaporation',
'Sunshine',
'WindGustSpeed',
'WindSpeed9am',
'WindSpeed3pm',
'Humidity9am',
'Humidity3pm',
'Pressure9am',
'Pressure3pm',
'Cloud9am',
'Cloud3pm',
'Temp9am',
'Temp3pm',
'Year',
'Month',
'Day']

```

Engineering missing values in numerical variables

```

# check missing values in numerical variables in X_train

X_train[numerical].isnull().sum()

```

MinTemp	495
MaxTemp	264
Rainfall	1139
Evaporation	48718
Sunshine	54314
WindGustSpeed	7367
WindSpeed9am	1086
WindSpeed3pm	2094
Humidity9am	1449
Humidity3pm	2890
Pressure9am	11212
Pressure3pm	11186
Cloud9am	43137
Cloud3pm	45768
Temp9am	740
Temp3pm	2171
Year	0
Month	0
Day	0

dtype: int64

check missing values in numerical variables in X_test

`X_test[numerical].isnull().sum()`

MinTemp	142
MaxTemp	58
Rainfall	267
Evaporation	12125
Sunshine	13502
WindGustSpeed	1903
WindSpeed9am	262
WindSpeed3pm	536
Humidity9am	325
Humidity3pm	720
Pressure9am	2802
Pressure3pm	2795
Cloud9am	10520
Cloud3pm	11326
Temp9am	164
Temp3pm	555
Year	0
Month	0
Day	0

dtype: int64

print percentage of missing values in the numerical variables in training set

`for col in numerical:`

```

if X_train[col].isnull().mean()>0:
    print(col, round(X_train[col].isnull().mean(),4))
MinTemp 0.0044
MaxTemp 0.0023
Rainfall 0.01
Evaporation 0.4283
Sunshine 0.4775
WindGustSpeed 0.0648
WindSpeed9am 0.0095
WindSpeed3pm 0.0184
Humidity9am 0.0127
Humidity3pm 0.0254
Pressure9am 0.0986
Pressure3pm 0.0983
Cloud9am 0.3792
Cloud3pm 0.4023
Temp9am 0.0065
Temp3pm 0.0191

```

Assumption

I assume that the data are missing completely at random (MCAR). There are two methods which can be used to impute missing values. One is mean or median imputation and other one is random sample imputation. When there are outliers in the dataset, we should use median imputation. So, I will use median imputation because median imputation is robust to outliers.

I will impute missing values with the appropriate statistical measures of the data, in this case median. Imputation should be done over the training set, and then propagated to the test set. It means that the statistical measures to be used to fill missing values both in train and test set, should be extracted from the train set only. This is to avoid overfitting.

```

# impute missing values in X_train and X_test with respective column
median in X_train

```

```

for df1 in [X_train, X_test]:
    for col in numerical:
        col_median=X_train[col].median()
        df1[col].fillna(col_median, inplace=True)

```

```

# check again missing values in numerical variables in X_train

```

```

X_train[numerical].isnull().sum()

```

```

MinTemp      0
MaxTemp      0
Rainfall     0
Evaporation   0
Sunshine     0

```

```
WindGustSpeed    0
WindSpeed9am     0
WindSpeed3pm     0
Humidity9am      0
Humidity3pm      0
Pressure9am      0
Pressure3pm      0
Cloud9am         0
Cloud3pm         0
Temp9am          0
Temp3pm          0
Year             0
Month            0
Day              0
dtype: int64
```

```
# check missing values in numerical variables in X_test
```

```
X_test[numerical].isnull().sum()
```

```
MinTemp          0
MaxTemp          0
Rainfall         0
Evaporation      0
Sunshine         0
WindGustSpeed    0
WindSpeed9am     0
WindSpeed3pm     0
Humidity9am      0
Humidity3pm      0
Pressure9am      0
Pressure3pm      0
Cloud9am         0
Cloud3pm         0
Temp9am          0
Temp3pm          0
Year             0
Month            0
Day              0
dtype: int64
```

Now, we can see that there are no missing values in the numerical columns of training and test set.

Engineering missing values in categorical variables

```
# print percentage of missing values in the categorical variables in training set
```

```
X_train[categorical].isnull().mean()
```

```
Location      0.000000
WindGustDir    0.065114
WindDir9am     0.070134
WindDir3pm     0.026443
RainToday      0.010013
dtype: float64
```

```
# print categorical variables with missing data
```

```
for col in categorical:
    if X_train[col].isnull().mean()>0:
        print(col, (X_train[col].isnull().mean()))
```

```
WindGustDir 0.06511419378659213
WindDir9am 0.07013379749283542
WindDir3pm 0.026443026179299188
RainToday 0.01001283471350458
```

```
# impute missing categorical variables with most frequent value
```

```
for df2 in [X_train, X_test]:
    df2['WindGustDir'].fillna(X_train['WindGustDir'].mode()[0],
inplace=True)
    df2['WindDir9am'].fillna(X_train['WindDir9am'].mode()[0],
inplace=True)
    df2['WindDir3pm'].fillna(X_train['WindDir3pm'].mode()[0],
inplace=True)
    df2['RainToday'].fillna(X_train['RainToday'].mode()[0],
inplace=True)
```

```
# check missing values in categorical variables in X_train
```

```
X_train[categorical].isnull().sum()
```

```
Location      0
WindGustDir    0
WindDir9am     0
WindDir3pm     0
RainToday      0
dtype: int64
```

```
# check missing values in categorical variables in X_test
```

```
X_test[categorical].isnull().sum()
```

```
Location      0
WindGustDir    0
WindDir9am     0
WindDir3pm     0
RainToday      0
dtype: int64
```


As a final check, I will check for missing values in X_train and X_test.

```
# check missing values in X_train
```

```
X_train.isnull().sum()
```

Location	0
MinTemp	0
MaxTemp	0
Rainfall	0
Evaporation	0
Sunshine	0
WindGustDir	0
WindGustSpeed	0
WindDir9am	0
WindDir3pm	0
WindSpeed9am	0
WindSpeed3pm	0
Humidity9am	0
Humidity3pm	0
Pressure9am	0
Pressure3pm	0
Cloud9am	0
Cloud3pm	0
Temp9am	0
Temp3pm	0
RainToday	0
Year	0
Month	0
Day	0

dtype: int64

```
# check missing values in X_test
```

```
X_test.isnull().sum()
```

Location	0
MinTemp	0
MaxTemp	0
Rainfall	0
Evaporation	0
Sunshine	0
WindGustDir	0
WindGustSpeed	0
WindDir9am	0
WindDir3pm	0
WindSpeed9am	0
WindSpeed3pm	0
Humidity9am	0
Humidity3pm	0

```

Pressure9am      0
Pressure3pm      0
Cloud9am         0
Cloud3pm         0
Temp9am          0
Temp3pm          0
RainToday        0
Year             0
Month            0
Day              0
dtype: int64

```

We can see that there are no missing values in X_train and X_test.

Engineering outliers in numerical variables

We have seen that the Rainfall, Evaporation, WindSpeed9am and WindSpeed3pm columns contain outliers. I will use top-coding approach to cap maximum values and remove outliers from the above variables.

```

def max_value(df3, variable, top):
    return np.where(df3[variable]>top, top, df3[variable])

for df3 in [X_train, X_test]:
    df3['Rainfall'] = max_value(df3, 'Rainfall', 3.2)
    df3['Evaporation'] = max_value(df3, 'Evaporation', 21.8)
    df3['WindSpeed9am'] = max_value(df3, 'WindSpeed9am', 55)
    df3['WindSpeed3pm'] = max_value(df3, 'WindSpeed3pm', 57)

```

```
X_train.Rainfall.max(), X_test.Rainfall.max()
```

```
(3.2, 3.2)
```

```
X_train.Evaporation.max(), X_test.Evaporation.max()
```

```
(21.8, 21.8)
```

```
X_train.WindSpeed9am.max(), X_test.WindSpeed9am.max()
```

```
(55.0, 55.0)
```

```
X_train.WindSpeed3pm.max(), X_test.WindSpeed3pm.max()
```

```
(57.0, 57.0)
```

```
X_train[numerical].describe()
```

	MinTemp	MaxTemp	Rainfall	Evaporation	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	12.193497	23.237216	0.675080	5.151606	
std	6.388279	7.094149	1.183837	2.823707	

min	-8.200000	-4.800000	0.000000	0.000000
25%	7.600000	18.000000	0.000000	4.000000
50%	12.000000	22.600000	0.000000	4.800000
75%	16.800000	28.200000	0.600000	5.400000
max	33.900000	48.100000	3.200000	21.800000

	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	8.041154	39.884074	13.978155	18.614756	
std	2.769480	13.116959	8.806558	8.685862	
min	0.000000	6.000000	0.000000	0.000000	
25%	8.200000	31.000000	7.000000	13.000000	
50%	8.500000	39.000000	13.000000	19.000000	
75%	8.700000	46.000000	19.000000	24.000000	
max	14.500000	135.000000	55.000000	57.000000	

	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	68.867486	51.509547	1017.640649	1015.241101	
std	18.935587	20.530723	6.738680	6.675168	
min	0.000000	0.000000	980.500000	977.100000	
25%	57.000000	37.000000	1013.500000	1011.000000	
50%	70.000000	52.000000	1017.600000	1015.200000	
75%	83.000000	65.000000	1021.800000	1019.400000	
max	100.000000	100.000000	1041.000000	1039.600000	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	4.651801	4.703588	16.995062	21.688643	
std	2.292726	2.117847	6.463772	6.855649	
min	0.000000	0.000000	-7.200000	-5.400000	
25%	3.000000	4.000000	12.300000	16.700000	
50%	5.000000	5.000000	16.700000	21.100000	
75%	6.000000	6.000000	21.500000	26.300000	
max	9.000000	8.000000	40.200000	46.700000	

	Year	Month	Day
count	113754.000000	113754.000000	113754.000000
mean	2012.759727	6.404021	15.710419
std	2.540419	3.427798	8.796821
min	2007.000000	1.000000	1.000000
25%	2011.000000	3.000000	8.000000
50%	2013.000000	6.000000	16.000000
75%	2015.000000	9.000000	23.000000
max	2017.000000	12.000000	31.000000

We can now see that the outliers in Rainfall, Evaporation, WindSpeed9am and WindSpeed3pm columns are capped.

Encode categorical variables

```
categorical
```

```
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday']
```

```
X_train[categorical].head()
```

	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday
110803	Witchcliffe	S	SSE	S	No
87289	Cairns	ENE	SSE	SE	Yes
134949	AliceSprings	E	NE	N	No
85553	Cairns	ESE	SSE	E	No
16110	Newcastle	W	N	SE	No

```
# encode RainToday variable
```

```
import category_encoders as ce
```

```
encoder = ce.BinaryEncoder(cols=['RainToday'])
```

```
X_train = encoder.fit_transform(X_train)
```

```
X_test = encoder.transform(X_test)
```

```
X_train.head()
```

	RainToday_0	RainToday_1	Location	MinTemp	MaxTemp
Rainfall \					
110803	0	1	Witchcliffe	13.9	22.6
0.2					
87289	1	0	Cairns	22.4	29.4
2.0					
134949	0	1	AliceSprings	9.7	36.2
0.0					
85553	0	1	Cairns	20.5	30.1
0.0					
16110	0	1	Newcastle	16.8	29.2
0.0					

	Evaporation	Sunshine	WindGustDir	WindGustSpeed	...
Humidity3pm \					
110803	4.8	8.5	S	41.0	...
55.0					
87289	6.0	6.3	ENE	33.0	...
59.0					
134949	11.4	12.3	E	31.0	...
2.0					
85553	8.8	11.1	ESE	37.0	...
53.0					
16110	4.8	8.5	W	39.0	...
53.0					

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm
Year \						
110803	1013.9	1013.4	5.0	5.0	18.8	20.4
2014						
87289	1016.9	1013.1	7.0	5.0	26.4	27.5
2015						
134949	1018.1	1013.6	1.0	1.0	28.5	35.0
2014						
85553	1014.1	1010.8	2.0	3.0	27.3	29.4
2010						
16110	1017.6	1015.2	5.0	8.0	22.2	27.0
2012						

	Month	Day
110803	4	25
87289	11	2
134949	10	19
85553	10	30
16110	11	8

[5 rows x 25 columns]

We can see that two additional variables `RainToday_0` and `RainToday_1` are created from `RainToday` variable.

Now, I will create the `X_train` training set.

```
X_train = pd.concat([X_train[numerical], X_train[['RainToday_0',
'RainToday_1']],
                    pd.get_dummies(X_train.Location),
                    pd.get_dummies(X_train.WindGustDir),
                    pd.get_dummies(X_train.WindDir9am),
                    pd.get_dummies(X_train.WindDir3pm)], axis=1)
```

`X_train.head()`

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine
WindGustSpeed \					
110803	13.9	22.6	0.2	4.8	8.5
41.0					
87289	22.4	29.4	2.0	6.0	6.3
33.0					
134949	9.7	36.2	0.0	11.4	12.3
31.0					
85553	20.5	30.1	0.0	8.8	11.1
37.0					
16110	16.8	29.2	0.0	4.8	8.5
39.0					

		WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	...	NNW
NW	S \						
110803		20.0	28.0	65.0	55.0	...	0
0	1						
87289		7.0	19.0	71.0	59.0	...	0
0	0						
134949		15.0	11.0	6.0	2.0	...	0
0	0						
85553		22.0	19.0	59.0	53.0	...	0
0	0						
16110		0.0	7.0	72.0	53.0	...	0
0	0						

	SE	SSE	SSW	SW	W	WNW	WSW
110803	0	0	0	0	0	0	0
87289	1	0	0	0	0	0	0
134949	0	0	0	0	0	0	0
85553	0	0	0	0	0	0	0
16110	1	0	0	0	0	0	0

[5 rows x 118 columns]

Similarly, I will create the X_test testing set.

```
X_test = pd.concat([X_test[numerical], X_test[['RainToday_0',
'RainToday_1']],
                    pd.get_dummies(X_test.Location),
                    pd.get_dummies(X_test.WindGustDir),
                    pd.get_dummies(X_test.WindDir9am),
                    pd.get_dummies(X_test.WindDir3pm)], axis=1)
```

X_test.head()

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine
WindGustSpeed \					
86232	17.4	29.0	0.0	3.6	11.1
33.0					
57576	6.8	14.4	0.8	0.8	8.5
46.0					
124071	10.1	15.4	3.2	4.8	8.5
31.0					
117955	14.4	33.4	0.0	8.0	11.6
41.0					
133468	6.8	14.3	3.2	0.2	7.3
28.0					

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	...	NNW	
NW	S \						
86232		11.0	19.0	63.0	61.0	...	0
0	0						

57576	17.0	22.0	80.0	55.0	...	0	
0 1							
124071	13.0	9.0	70.0	61.0	...	0	
0 0							
117955	9.0	17.0	40.0	23.0	...	0	
0 0							
133468	15.0	13.0	92.0	47.0	...	0	
0 0							
	SE	SSE	SSW	SW	W	WNW	WSW
86232	0	0	0	0	0	0	0
57576	0	0	0	0	0	0	0
124071	0	1	0	0	0	0	0
117955	0	0	0	1	0	0	0
133468	0	0	0	0	0	0	0
[5 rows x 118 columns]							

We now have training and testing set ready for model building. Before that, we should map all the feature variables onto the same scale. It is called **feature scaling**. I will do it as follows.

11. Feature Scaling

X_train.describe()

	MinTemp	MaxTemp	Rainfall	Evaporation	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	12.193497	23.237216	0.675080	5.151606	
std	6.388279	7.094149	1.183837	2.823707	
min	-8.200000	-4.800000	0.000000	0.000000	
25%	7.600000	18.000000	0.000000	4.000000	
50%	12.000000	22.600000	0.000000	4.800000	
75%	16.800000	28.200000	0.600000	5.400000	
max	33.900000	48.100000	3.200000	21.800000	
	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	8.041154	39.884074	13.978155	18.614756	
std	2.769480	13.116959	8.806558	8.685862	
min	0.000000	6.000000	0.000000	0.000000	
25%	8.200000	31.000000	7.000000	13.000000	
50%	8.500000	39.000000	13.000000	19.000000	
75%	8.700000	46.000000	19.000000	24.000000	
max	14.500000	135.000000	55.000000	57.000000	
	Humidity9am	Humidity3pm	...	NNW	\
count	113754.000000	113754.000000	...	113754.000000	
mean	68.867486	51.509547	...	0.054530	
std	18.935587	20.530723	...	0.227061	
min	0.000000	0.000000	...	0.000000	

25%	57.000000	37.000000	...	0.000000
50%	70.000000	52.000000	...	0.000000
75%	83.000000	65.000000	...	0.000000
max	100.000000	100.000000	...	1.000000

	NW	S	SE	SSE \
count	113754.000000	113754.000000	113754.000000	113754.000000
mean	0.060288	0.067259	0.101605	0.064059
std	0.238021	0.250471	0.302130	0.244860
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	SSW	SW	W	WNW \
count	113754.000000	113754.000000	113754.000000	113754.000000
mean	0.056402	0.064464	0.069334	0.060798
std	0.230698	0.245578	0.254022	0.238960
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	WSW
count	113754.000000
mean	0.065483
std	0.247378
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 118 columns]

cols = X_train.columns

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

X_train = pd.DataFrame(X_train, columns=[cols])

X_test = pd.DataFrame(X_test, columns=[cols])

X_train.describe()

	MinTemp	MaxTemp	Rainfall	Evaporation	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	0.484406	0.530004	0.210962	0.236312	
std	0.151741	0.134105	0.369949	0.129528	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.375297	0.431002	0.000000	0.183486	
50%	0.479810	0.517958	0.000000	0.220183	
75%	0.593824	0.623819	0.187500	0.247706	
max	1.000000	1.000000	1.000000	1.000000	

	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	0.554562	0.262667	0.254148	0.326575	
std	0.190999	0.101682	0.160119	0.152384	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.565517	0.193798	0.127273	0.228070	
50%	0.586207	0.255814	0.236364	0.333333	
75%	0.600000	0.310078	0.345455	0.421053	
max	1.000000	1.000000	1.000000	1.000000	

	Humidity9am	Humidity3pm	...	NNW	\
count	113754.000000	113754.000000	...	113754.000000	
mean	0.688675	0.515095	...	0.054530	
std	0.189356	0.205307	...	0.227061	
min	0.000000	0.000000	...	0.000000	
25%	0.570000	0.370000	...	0.000000	
50%	0.700000	0.520000	...	0.000000	
75%	0.830000	0.650000	...	0.000000	
max	1.000000	1.000000	...	1.000000	

	NW	S	SE	SSE	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	0.060288	0.067259	0.101605	0.064059	
std	0.238021	0.250471	0.302130	0.244860	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	SSW	SW	W	WNW	\
count	113754.000000	113754.000000	113754.000000	113754.000000	
mean	0.056402	0.064464	0.069334	0.060798	
std	0.230698	0.245578	0.254022	0.238960	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	

max	1.000000	1.000000	1.000000	1.000000
-----	----------	----------	----------	----------

	WSW
count	113754.000000
mean	0.065483
std	0.247378
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 118 columns]

We now have `X_train` dataset ready to be fed into the Logistic Regression classifier. I will do it as follows.

12. Model training

```
# train a logistic regression model on the training set
from sklearn.linear_model import LogisticRegression

# instantiate the model
logreg = LogisticRegression(solver='liblinear', random_state=0)

# fit the model
logreg.fit(X_train, y_train)

LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=0,
solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False)
```

13. Predict results

```
y_pred_test = logreg.predict(X_test)

y_pred_test
array(['No', 'No', 'No', ..., 'No', 'No', 'Yes'], dtype=object)
```

predict_proba method

predict_proba method gives the probabilities for the target variable(0 and 1) in this case, in array form.

0 is for probability of no rain and 1 is for probability of rain.

```
# probability of getting output as 0 - no rain
logreg.predict_proba(X_test)[: ,0]
array([0.91387232, 0.83563172, 0.82035588, ..., 0.97674036,
       0.7985333 ,
       0.3073458 ])

# probability of getting output as 1 - rain
logreg.predict_proba(X_test)[: ,1]
array([0.08612768, 0.16436828, 0.17964412, ..., 0.02325964,
       0.2014667 ,
       0.6926542  ])
```

14. Check accuracy score

```
from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test,
y_pred_test)))

Model accuracy score: 0.8501
```

Here, **y_test** are the true class labels and **y_pred_test** are the predicted class labels in the test-set.

Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
y_pred_train = logreg.predict(X_train)

y_pred_train
array(['No', 'No', 'No', ..., 'No', 'No', 'No'], dtype=object)

print('Training-set accuracy score: {0:0.4f}'.
format(accuracy_score(y_train, y_pred_train)))

Training-set accuracy score: 0.8476
```

Check for overfitting and underfitting

```
# print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg.score(X_train,
y_train)))
```

```
print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))
```

Training set score: 0.8476

Test set score: 0.8501

The training-set accuracy score is 0.8476 while the test-set accuracy to be 0.8501. These two values are quite comparable. So, there is no question of overfitting.

In Logistic Regression, we use default value of $C = 1$. It provides good performance with approximately 85% accuracy on both the training and the test set. But the model performance on both the training and test set are very comparable. It is likely the case of underfitting.

I will increase C and fit a more flexible model.

```
# fit the Logsitic Regression model with C=100
```

```
# instantiate the model
```

```
logreg100 = LogisticRegression(C=100, solver='liblinear',  
                                random_state=0)
```

```
# fit the model
```

```
logreg100.fit(X_train, y_train)
```

```
LogisticRegression(C=100, class_weight=None, dual=False,  
fit_intercept=True,  
                    intercept_scaling=1, max_iter=100, multi_class='warn',  
                    n_jobs=None, penalty='l2', random_state=0,  
solver='liblinear',  
                    tol=0.0001, verbose=0, warm_start=False)
```

```
# print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(logreg100.score(X_train,  
y_train)))
```

```
print('Test set score: {:.4f}'.format(logreg100.score(X_test,  
y_test)))
```

Training set score: 0.8478

Test set score: 0.8505

We can see that, $C=100$ results in higher test set accuracy and also a slightly increased training set accuracy. So, we can conclude that a more complex model should perform better.

Now, I will investigate, what happens if we use more regularized model than the default value of $C=1$, by setting $C=0.01$.

```
# fit the Logsitic Regression model with C=001
```

```

# instantiate the model
logreg001 = LogisticRegression(C=0.01, solver='liblinear',
random_state=0)

# fit the model
logreg001.fit(X_train, y_train)

LogisticRegression(C=0.01, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=0,
solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False)

# print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg001.score(X_train,
y_train)))

print('Test set score: {:.4f}'.format(logreg001.score(X_test,
y_test)))

Training set score: 0.8409
Test set score: 0.8448

```

So, if we use more regularized model by setting $C=0.01$, then both the training and test set accuracy decrease relative to the default parameters.

Compare model accuracy with null accuracy

So, the model accuracy is 0.8501. But, we cannot say that our model is very good based on the above accuracy. We must compare it with the **null accuracy**. Null accuracy is the accuracy that could be achieved by always predicting the most frequent class.

So, we should first check the class distribution in the test set.

```

# check class distribution in test set

y_test.value_counts()

No      22067
Yes      6372
Name: RainTomorrow, dtype: int64

```

We can see that the occurrences of most frequent class is 22067. So, we can calculate null accuracy by dividing 22067 by total number of occurrences.

```

# check null accuracy score

```

```
null_accuracy = (22067/(22067+6372))  
print('Null accuracy score: {0:0.4f}'.format(null_accuracy))  
Null accuracy score: 0.7759
```

We can see that our model accuracy score is 0.8501 but null accuracy score is 0.7759. So, we can conclude that our Logistic Regression model is doing a very good job in predicting the class labels.

Now, based on the above analysis we can conclude that our classification model accuracy is very good. Our model is doing a very good job in terms of predicting the class labels.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called **Confusion matrix** that comes to our rescue.

15. Confusion matrix

A confusion matrix is a tool for summarizing the performance of a classification algorithm. A confusion matrix will give us a clear picture of classification model performance and the types of errors produced by the model. It gives us a summary of correct and incorrect predictions broken down by each category. The summary is represented in a tabular form.

Four types of outcomes are possible while evaluating a classification model performance. These four outcomes are described below:-

True Positives (TP) – True Positives occur when we predict an observation belongs to a certain class and the observation actually belongs to that class.

True Negatives (TN) – True Negatives occur when we predict an observation does not belong to a certain class and the observation actually does not belong to that class.

False Positives (FP) – False Positives occur when we predict an observation belongs to a certain class but the observation actually does not belong to that class. This type of error is called **Type I error**.

False Negatives (FN) – False Negatives occur when we predict an observation does not belong to a certain class but the observation actually belongs to that class. This is a very serious error and it is called **Type II error**.

These four outcomes are summarized in a confusion matrix given below.

```
# Print the Confusion Matrix and slice it into four pieces  
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred_test)  
print('Confusion matrix\n\n', cm)
```

```

print('\nTrue Positives(TP) = ', cm[0,0])
print('\nTrue Negatives(TN) = ', cm[1,1])
print('\nFalse Positives(FP) = ', cm[0,1])
print('\nFalse Negatives(FN) = ', cm[1,0])

```

Confusion matrix

```

[[20892  1175]
 [ 3087  3285]]

```

True Positives(TP) = 20892

True Negatives(TN) = 3285

False Positives(FP) = 1175

False Negatives(FN) = 3087

The confusion matrix shows $20892 + 3285 = 24177$ correct predictions and $3087 + 1175 = 4262$ incorrect predictions.

In this case, we have

- True Positives (Actual Positive:1 and Predict Positive:1) - 20892
- True Negatives (Actual Negative:0 and Predict Negative:0) - 3285
- False Positives (Actual Negative:0 but Predict Positive:1) - 1175 (Type I error)
- False Negatives (Actual Positive:1 but Predict Negative:0) - 3087 (Type II error)


visualize confusion matrix with seaborn heatmap

```

cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1',
                                           'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict
Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
<matplotlib.axes._subplots.AxesSubplot at 0xacc3104f60>

```



16. Classification metrics

Classification Report

Classification report is another way to evaluate the classification model performance. It displays the **precision**, **recall**, **f1** and **support** scores for the model. I have described these terms in later.

We can print a classification report as follows:-

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
No	0.87	0.95	0.91	22067
Yes	0.74	0.52	0.61	6372
micro avg	0.85	0.85	0.85	28439
macro avg	0.80	0.73	0.76	28439
weighted avg	0.84	0.85	0.84	28439

Classification accuracy

```
TP = cm[0,0]
TN = cm[1,1]
```



```

FP = cm[0,1]
FN = cm[1,0]

# print classification accuracy

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy :
{0:0.4f}'.format(classification_accuracy))

Classification accuracy : 0.8501

```

Classification error

```

# print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))

Classification error : 0.1499

```

Precision

Precision can be defined as the percentage of correctly predicted positive outcomes out of all the predicted positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true and false positives (TP + FP).

So, **Precision** identifies the proportion of correctly predicted positive outcome. It is more concerned with the positive class than the negative class.

Mathematically, precision can be defined as the ratio of TP to (TP + FP).

```

# print precision score

precision = TP / float(TP + FP)

print('Precision : {0:0.4f}'.format(precision))

Precision : 0.9468

```

Recall

Recall can be defined as the percentage of correctly predicted positive outcomes out of all the actual positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true positives and false negatives (TP + FN). **Recall** is also called **Sensitivity**.

Recall identifies the proportion of correctly predicted actual positives.

Mathematically, recall can be given as the ratio of TP to (TP + FN).

```
recall = TP / float(TP + FN)
print('Recall or Sensitivity : {0:0.4f}'.format(recall))
Recall or Sensitivity : 0.8713
```

True Positive Rate

True Positive Rate is synonymous with **Recall**.

```
true_positive_rate = TP / float(TP + FN)
print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
True Positive Rate : 0.8713
```

False Positive Rate

```
false_positive_rate = FP / float(FP + TN)
print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))
False Positive Rate : 0.2635
```

Specificity

```
specificity = TN / (TN + FP)
print('Specificity : {0:0.4f}'.format(specificity))
Specificity : 0.7365
```

f1-score

f1-score is the weighted harmonic mean of precision and recall. The best possible **f1-score** would be 1.0 and the worst would be 0.0. **f1-score** is the harmonic mean of precision and recall. So, **f1-score** is always lower than accuracy measures as they embed precision and recall into their computation. The weighted average of **f1-score** should be used to compare classifier models, not global accuracy.

Support

Support is the actual number of occurrences of the class in our dataset.

17. Adjusting the threshold level

```
# print the first 10 predicted probabilities of two classes- 0 and 1
```

```
y_pred_prob = logreg.predict_proba(X_test)[0:10]
```

```
y_pred_prob
```

```
array([[0.91387232, 0.08612768],
       [0.83563172, 0.16436828],
       [0.82035588, 0.17964412],
       [0.99025882, 0.00974118],
       [0.95726809, 0.04273191],
       [0.97994232, 0.02005768],
       [0.17838588, 0.82161412],
       [0.23482434, 0.76517566],
       [0.90050811, 0.09949189],
       [0.85480088, 0.14519912]])
```

Observations

- In each row, the numbers sum to 1.
- There are 2 columns which correspond to 2 classes - 0 and 1.
 - Class 0 - predicted probability that there is no rain tomorrow.
 - Class 1 - predicted probability that there is rain tomorrow.
- Importance of predicted probabilities
 - We can rank the observations by probability of rain or no rain.
- predict_proba process
 - Predicts the probabilities
 - Choose the class with the highest probability
- Classification threshold level
 - There is a classification threshold level of 0.5.
 - Class 1 - probability of rain is predicted if probability > 0.5.
 - Class 0 - probability of no rain is predicted if probability < 0.5.

```
# store the probabilities in dataframe
```

```
y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of - No  
rain tomorrow (0)', 'Prob of - Rain tomorrow (1)'])
```

```
y_pred_prob_df
```

	Prob of - No rain tomorrow (0)	Prob of - Rain tomorrow (1)
0	0.913872	0.086128

1	0.835632	0.164368
2	0.820356	0.179644
3	0.990259	0.009741
4	0.957268	0.042732
5	0.979942	0.020058
6	0.178386	0.821614
7	0.234824	0.765176
8	0.900508	0.099492
9	0.854801	0.145199

```
# print the first 10 predicted probabilities for class 1 - Probability of rain
```

```
logreg.predict_proba(X_test)[0:10, 1]
```

```
array([0.08612768, 0.16436828, 0.17964412, 0.00974118, 0.04273191,
       0.02005768, 0.82161412, 0.76517566, 0.09949189, 0.14519912])
```

```
# store the predicted probabilities for class 1 - Probability of rain
```

```
y_pred1 = logreg.predict_proba(X_test)[: , 1]
```

```
# plot histogram of predicted probabilities
```

```
# adjust the font size
```

```
plt.rcParams['font.size'] = 12
```

```
# plot histogram with 10 bins
```

```
plt.hist(y_pred1, bins = 10)
```

```
# set the title of predicted probabilities
```

```
plt.title('Histogram of predicted probabilities of rain')
```

```
# set the x-axis limit
```


```
plt.xlim(0,1)
```

```
# set the title
```

```
plt.xlabel('Predicted probabilities of rain')
```

```
plt.ylabel('Frequency')
```

```
Text(0,0.5,'Frequency')
```



Observations

- We can see that the above histogram is highly positive skewed.
- The first column tell us that there are approximately 15000 observations with probability between 0.0 and 0.1.
- There are small number of observations with probability > 0.5 .
- So, these small number of observations predict that there will be rain tomorrow.
- Majority of observations predict that there will be no rain tomorrow.

Lower the threshold

```
from sklearn.preprocessing import binarize
for i in range(1,5):
    cm1=0
    y_pred1 = logreg.predict_proba(X_test)[:,-1]
    y_pred1 = y_pred1.reshape(-1,1)
    y_pred2 = binarize(y_pred1, i/10)
    y_pred2 = np.where(y_pred2 == 1, 'Yes', 'No')
```

```

cm1 = confusion_matrix(y_test, y_pred2)

print ('With',i/10,'threshold the Confusion Matrix is ', '\n\
n',cm1, '\n\n',

      'with',cm1[0,0]+cm1[1,1], 'correct predictions, ', '\n\n',
      cm1[0,1], 'Type I errors( False Positives), ', '\n\n',
      cm1[1,0], 'Type II errors( False Negatives), ', '\n\n',
      'Accuracy score: ', (accuracy_score(y_test, y_pred2)), '\n\
n',

      'Sensitivity: ',cm1[1,1]/(float(cm1[1,1]+cm1[1,0])), '\n\
n',

      'Specificity: ',cm1[0,0]/(float(cm1[0,0]+cm1[0,1])), '\n\n',
      '===== ', '\
n\n')

```

With 0.1 threshold the Confusion Matrix is

```

[[12726  9341]
 [  547  5825]]

```

with 18551 correct predictions,

9341 Type I errors(False Positives),

547 Type II errors(False Negatives),

Accuracy score: 0.6523084496641935

Sensitivity: 0.9141556811048337

Specificity: 0.5766982371867494

=====

With 0.2 threshold the Confusion Matrix is

```

[[17067  5000]
 [ 1233  5139]]

```

with 22206 correct predictions,

5000 Type I errors(False Positives),

1233 Type II errors(False Negatives),

Accuracy score: 0.7808291430781673

Sensitivity: 0.806497175141243

Specificity: 0.7734173199800607

=====

With 0.3 threshold the Confusion Matrix is

```
[[19080  2987]
 [ 1873  4499]]
```

with 23579 correct predictions,

2987 Type I errors(False Positives),

1873 Type II errors(False Negatives),

Accuracy score: 0.8291079151868912

Sensitivity: 0.7060577526679221

Specificity: 0.8646395069560883

=====

With 0.4 threshold the Confusion Matrix is

```
[[20191  1876]
 [ 2517  3855]]
```

with 24046 correct predictions,

1876 Type I errors(False Positives),

2517 Type II errors(False Negatives),

Accuracy score: 0.845529027040332

Sensitivity: 0.6049905838041432

Specificity: 0.9149861784565188

=====

Comments

- In binary problems, the threshold of 0.5 is used by default to convert predicted probabilities into class predictions.
- Threshold can be adjusted to increase sensitivity or specificity.
- Sensitivity and specificity have an inverse relationship. Increasing one would always decrease the other and vice versa.
- We can see that increasing the threshold level results in increased accuracy.
- Adjusting the threshold level should be one of the last step you do in the model-building process.

18. ROC - AUC

ROC Curve

Another tool to measure the classification model performance visually is **ROC Curve**. ROC Curve stands for **Receiver Operating Characteristic Curve**. An **ROC Curve** is a plot which shows the performance of a classification model at various classification threshold levels.

The **ROC Curve** plots the **True Positive Rate (TPR)** against the **False Positive Rate (FPR)** at various threshold levels.

True Positive Rate (TPR) is also called **Recall**. It is defined as the ratio of TP to (TP + FN) .

False Positive Rate (FPR) is defined as the ratio of FP to (FP + TN) .

In the ROC Curve, we will focus on the TPR (True Positive Rate) and FPR (False Positive Rate) of a single point. This will give us the general performance of the ROC curve which consists of the TPR and FPR at various threshold levels. So, an ROC Curve plots TPR vs FPR at different classification threshold levels. If we lower the threshold levels, it may result in more items being classified as positive. It will increase both True Positives (TP) and False Positives (FP).

```
# plot ROC Curve

from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = 'Yes')

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)


plt.plot([0,1], [0,1], 'k--' )

plt.rcParams['font.size'] = 12

plt.title('ROC curve for RainTomorrow classifier')
```



```
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
```



ROC curve help us to choose a threshold level that balances sensitivity and specificity for a particular context.

ROC AUC

ROC AUC stands for **Receiver Operating Characteristic - Area Under Curve**. It is a technique to compare classifier performance. In this technique, we measure the **area under the curve (AUC)**. A perfect classifier will have a ROC AUC equal to 1, whereas a purely random classifier will have a ROC AUC equal to 0.5.

So, **ROC AUC** is the percentage of the ROC plot that is underneath the curve.

```
# compute ROC AUC

from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC AUC : 0.8729

Comments

- ROC AUC is a single number summary of classifier performance. The higher the value, the better the classifier.
- ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a good job in predicting whether it will rain tomorrow or not.

```
# calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(logreg, X_train, y_train,
cv=5, scoring='roc_auc').mean()

print('Cross validated ROC AUC :
{:.4f}'.format(Cross_validated_ROC_AUC))

Cross validated ROC AUC : 0.8695
```

Model evaluation and improvement

In this section, I will employ several techniques to improve the model performance. I will discuss 3 techniques which are used in practice for performance improvement. These are recursive feature elimination, k-fold cross validation and hyperparameter optimization using GridSearchCV.

19. Recursive Feature Elimination with Cross Validation

Recursive feature elimination (RFE) is a feature selection technique that helps us to select best features from the given number of features. At first, the model is built on all the given features. Then, it removes the least useful predictor and build the model again. This process is repeated until all the unimportant features are removed from the model.

Recursive Feature Elimination with Cross-Validated (RFECV) feature selection technique selects the best subset of features for the estimator by removing 0 to N features iteratively using recursive feature elimination. Then it selects the best subset based on the accuracy or cross-validation score or roc-auc of the model. Recursive feature elimination technique eliminates n features from a model by fitting the model multiple times and at each step, removing the weakest features.

I will use this technique to select best features from this model.

```
from sklearn.feature_selection import RFECV

rfecv = RFECV(estimator=logreg, step=1, cv=5, scoring='accuracy')
```

```

rfecv = rfecv.fit(X_train, y_train)
print("Optimal number of features : %d" % rfecv.n_features_)
Optimal number of features : 112
# transform the training data
X_train_rfecv = rfecv.transform(X_train)

# train classifier
logreg.fit(X_train_rfecv, y_train)
LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=0,
solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False)

# test classifier on test data
X_test_rfecv = rfecv.transform(X_test)
y_pred_rfecv = logreg.predict(X_test_rfecv)
# print mean accuracy on transformed test data and labels
print ("Classifier score:
{:.4f}".format(logreg.score(X_test_rfecv,y_test)))
Classifier score: 0.8500

```

Our original model accuracy score is 0.8501 whereas accuracy score after RFECV is 0.8500. So, we can obtain approximately similar accuracy but with reduced or optimal set of features.

Confusion-matrix revisited

I will again plot the confusion-matrix for this model to get an idea of errors our model is making.

```

from sklearn.metrics import confusion_matrix
cm1 = confusion_matrix(y_test, y_pred_rfecv)
print('Confusion matrix\n\n', cm1)
print('\nTrue Positives(TP1) = ', cm1[0,0])
print('\nTrue Negatives(TN1) = ', cm1[1,1])

```

```
print('\nFalse Positives(FP1) = ', cm1[0,1])
```

```
print('\nFalse Negatives(FN1) = ', cm1[1,0])
```

Confusion matrix

```
[[20893  1174]
 [ 3091  3281]]
```

True Positives(TP1) = 20893

True Negatives(TN1) = 3281

False Positives(FP1) = 1174

False Negatives(FN1) = 3091

We can see that in the original model, we have FP = 1175 whereas FP1 = 1174. So, we get approximately same number of false positives. Also, FN = 3087 whereas FN1 = 3091. So, we get slightly higher false negatives.

20. k-Fold Cross Validation

Applying 10-Fold Cross Validation

```
from sklearn.model_selection import cross_val_score
```

```
scores = cross_val_score(logreg, X_train, y_train, cv = 5,
scoring='accuracy')
```

```
print('Cross-validation scores:{}'.format(scores))
```

```
Cross-validation scores:[0.84690783 0.84624852 0.84633642 0.84958903
0.84773626]
```

We can summarize the cross-validation accuracy by calculating its mean.

compute Average cross-validation score

```
print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

```
Average cross-validation score: 0.8474
```

Our, original model score is found to be 0.8476. The average cross-validation score is 0.8474. So, we can conclude that cross-validation does not result in performance improvement.

21. Hyperparameter Optimization using GridSearch CV

```
from sklearn.model_selection import GridSearchCV

parameters = [{'penalty':['l1','l2']},
               {'C':[1, 10, 100, 1000]}]

grid_search = GridSearchCV(estimator = logreg,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 5,
                           verbose=0)

grid_search.fit(X_train, y_train)

GridSearchCV(cv=5, error_score='raise-deprecating',
             estimator=LogisticRegression(C=1.0, class_weight=None,
             dual=False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='warn',
             n_jobs=None, penalty='l2', random_state=0,
             solver='liblinear',
             tol=0.0001, verbose=0, warm_start=False),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid=[{'penalty': ['l1', 'l2']}, {'C': [1, 10, 100,
             1000]}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring='accuracy', verbose=0)

# examine the best model

# best score achieved during the GridSearchCV
print('GridSearch CV best score : {:.4f}\n\n'
      .format(grid_search.best_score_))

# print parameters that give the best results
print('Parameters that give the best results :','\n\n',
      (grid_search.best_params_))

# print estimator that was chosen by the GridSearch
print('\n\nEstimator that was chosen by the search :','\n\n',
      (grid_search.best_estimator_))

GridSearch CV best score : 0.8474

Parameters that give the best results :
```

```
{'penalty': 'l1'}
```

Estimator that was chosen by the search :

```
LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l1', random_state=0,
solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False)
```

```
# calculate GridSearch CV score on test set
```

```
print('GridSearch CV score on test set:
{0:0.4f}'.format(grid_search.score(X_test, y_test)))
```

```
GridSearch CV score on test set: 0.8507
```

Comments

- Our original model test accuracy is 0.8501 while GridSearch CV accuracy is 0.8507.
- We can see that GridSearch CV improve the performance for this particular model.

22. Results and Conclusion

1. The logistic regression model accuracy score is 0.8501. So, the model does a very good job in predicting whether or not it will rain tomorrow in Australia.
2. Small number of observations predict that there will be rain tomorrow. Majority of observations predict that there will be no rain tomorrow.
3. The model shows no signs of overfitting.
4. Increasing the value of C results in higher test set accuracy and also a slightly increased training set accuracy. So, we can conclude that a more complex model should perform better.
5. Increasing the threshold level results in increased accuracy.
6. ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a good job in predicting whether it will rain tomorrow or not.
7. Our original model accuracy score is 0.8501 whereas accuracy score after RFECV is 0.8500. So, we can obtain approximately similar accuracy but with reduced set of features.
8. In the original model, we have FP = 1175 whereas FP1 = 1174. So, we get approximately same number of false positives. Also, FN = 3087 whereas FN1 = 3091. So, we get slightly higher false negatives.

9. Our, original model score is found to be 0.8476. The average cross-validation score is 0.8474. So, we can conclude that cross-validation does not result in performance improvement.
10. Our original model test accuracy is 0.8501 while GridSearch CV accuracy is 0.8507. We can see that GridSearch CV improve the performance for this particular model.