

# Naive Bayes Classification with Python and Scikit-Learn

In this project, I implement Naive Bayes Classification algorithm with Python and Scikit-Learn. I build a Naive Bayes Classifier to predict whether a person makes over 50K a year. I have used the **Adult Data Set** for this project. I have downloaded this dataset from the UCI Machine Learning Repository website.

## Table of Contents

1. Introduction to Naive Bayes Classification algorithm
2. Naive Bayes algorithm intuition
3. The problem statement
4. Dataset description
5. Import libraries
6. Import dataset
7. Exploratory data analysis
8. Declare feature vector and target variable
9. Split data into separate training and test set
10. Feature engineering
11. Feature scaling
12. Model training
13. Predict the test-set results
14. Check the accuracy score
15. Confusion matrix
16. Classification metrics
17. Calculate class probabilities
18. ROC - AUC
19. k-Fold Cross Validation
20. Results and conclusion

## 1. Introduction to Naive Bayes Classification algorithm

In machine learning, Naïve Bayes classification is a straightforward and powerful algorithm for the classification task. Naïve Bayes classification is based on applying Bayes' theorem with strong independence assumption between the features. Naïve Bayes classification produces good results when we use it for textual data analysis such as Natural Language Processing.

Naïve Bayes models are also known as **simple Bayes** or **independent Bayes**. All these names refer to the application of Bayes' theorem in the classifier's decision rule. Naïve Bayes classifier applies the Bayes' theorem in practice. This classifier brings the power of Bayes' theorem to machine learning.

## 2. Naive Bayes algorithm intuition

Naïve Bayes Classifier uses the Bayes' theorem to predict membership probabilities for each class such as the probability that given record or data point belongs to a particular class. The class with the highest probability is considered as the most likely class. This is also known as the **Maximum A Posteriori (MAP)**.

The **MAP for a hypothesis with 2 events A and B** is

**MAP (A)**

$$= \max (P (A | B))$$

$$= \max (P (B | A) * P (A))/P (B)$$

$$= \max (P (B | A) * P (A))$$

Here,  $P (B)$  is evidence probability. It is used to normalize the result. It remains the same, So, removing it would not affect the result.

Naïve Bayes Classifier assumes that all the features are unrelated to each other. Presence or absence of a feature does not influence the presence or absence of any other feature.

In real world datasets, we test a hypothesis given multiple evidence on features. So, the calculations become quite complicated. To simplify the work, the feature independence approach is used to uncouple multiple evidence and treat each as an independent one.

## 3. The problem statement

In this project, I try to make predictions where the prediction task is to determine whether a person makes over 50K a year. I implement Naive Bayes Classification with Python and Scikit-Learn. So, to answer the question, I build a Naive Bayes classifier to predict whether a person makes over 50K a year.

## 4. Dataset description

I have used the **Adult Data Set** for this project. I have downloaded this dataset from the UCI Machine Learning Repository website. The data set can be found at the following url:-

<https://archive.ics.uci.edu/ml/datasets/Adult>

## 5. Import libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
import warnings

warnings.filterwarnings('ignore')
```

## 6. Import dataset

```
data = 'C:/datasets/adult.data'

df = pd.read_csv(data, header=None, sep=',\s')
```

## 7. Exploratory data analysis

Now, I will explore the data to gain insights about the data.

```
# view dimensions of dataset

df.shape

(32561, 15)
```

We can see that there are 32561 instances and 15 attributes in the data set.

### View top 5 rows of dataset

```
# preview the dataset

df.head()
```

	0	1	2	3	4	5	\
0	39	State-gov	77516	Bachelors	13	Never-married	
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	
2	38	Private	215646	HS-grad	9	Divorced	
3	53	Private	234721	11th	7	Married-civ-spouse	
4	28	Private	338409	Bachelors	13	Married-civ-spouse	

	6	7	8	9	10	11	12	\
0	Adm-clerical	Not-in-family	White	Male	2174	0	40	
1	Exec-managerial	Husband	White	Male	0	0	13	
2	Handlers-cleaners	Not-in-family	White	Male	0	0	40	
3	Handlers-cleaners	Husband	Black	Male	0	0	40	
4	Prof-specialty	Wife	Black	Female	0	0	40	

	13	14
0	United-States	<=50K
1	United-States	<=50K
2	United-States	<=50K
3	United-States	<=50K
4	Cuba	<=50K

## Rename column names

We can see that the dataset does not have proper column names. The columns are merely labelled as 0,1,2.... and so on. We should give proper names to the columns. I will do it as follows:-

```
col_names = ['age', 'workclass', 'fnlwgt', 'education',
             'education_num', 'marital_status', 'occupation', 'relationship',
             'race', 'sex', 'capital_gain', 'capital_loss',
             'hours_per_week', 'native_country', 'income']

df.columns = col_names

df.columns
Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
      'marital_status', 'occupation', 'relationship', 'race', 'sex',
      'capital_gain', 'capital_loss', 'hours_per_week',
      'native_country',
      'income'],
      dtype='object')
```

*# let's again preview the dataset*

```
df.head()
```

	age	workclass	fnlwgt	education	education_num \	
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital_status	occupation	relationship	race	sex
0	Never-married	Adm-clerical	Not-in-family	White	Male
1	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

	capital_gain	capital_loss	hours_per_week	native_country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K

3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

We can see that the column names are renamed. Now, the columns have meaningful names.

## View summary of dataset

```
# view summary of dataset

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
age                32561 non-null int64
workclass          32561 non-null object
fnlwgt             32561 non-null int64
education          32561 non-null object
education_num      32561 non-null int64
marital_status     32561 non-null object
occupation         32561 non-null object
relationship       32561 non-null object
race               32561 non-null object
sex                32561 non-null object
capital_gain       32561 non-null int64
capital_loss       32561 non-null int64
hours_per_week     32561 non-null int64
native_country     32561 non-null object
income             32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

We can see that there are no missing values in the dataset. I will confirm this further.

## Types of variables

In this section, I segregate the dataset into categorical and numerical variables. There are a mixture of categorical and numerical variables in the dataset. Categorical variables have data type object. Numerical variables have data type int64.

First of all, I will explore categorical variables.

## Explore categorical variables

```
# find categorical variables

categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :\n\n', categorical)
```

There are 9 categorical variables

The categorical variables are :

```
['workclass', 'education', 'marital_status', 'occupation',  
'relationship', 'race', 'sex', 'native_country', 'income']
```

```
# view the categorical variables
```

```
df[categorical].head()
```

	workclass	education	marital_status	occupation
0	State-gov	Bachelors	Never-married	Adm-clerical
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial
2	Private	HS-grad	Divorced	Handlers-cleaners
3	Private	11th	Married-civ-spouse	Handlers-cleaners
4	Private	Bachelors	Married-civ-spouse	Prof-specialty

	relationship	race	sex	native_country	income
0	Not-in-family	White	Male	United-States	<=50K
1	Husband	White	Male	United-States	<=50K
2	Not-in-family	White	Male	United-States	<=50K
3	Husband	Black	Male	United-States	<=50K
4	Wife	Black	Female	Cuba	<=50K

## Summary of categorical variables

- There are 9 categorical variables.
- The categorical variables are given by workclass, education, marital\_status, occupation, relationship, race, sex, native\_country and income.
- income is the target variable.

## Explore problems within categorical variables

First, I will explore the categorical variables.

## Missing values in categorical variables

```
# check missing values in categorical variables
```

```
df[categorical].isnull().sum()
```

workclass	0
education	0

```
marital_status    0
occupation        0
relationship      0
race              0
sex               0
native_country    0
income            0
dtype: int64
```

We can see that there are no missing values in the categorical variables. I will confirm this further.

## Frequency counts of categorical variables

Now, I will check the frequency counts of categorical variables.

```
# view frequency counts of values in categorical variables
```

```
for var in categorical:
```

```
    print(df[var].value_counts())
```

```
Private          22696
Self-emp-not-inc  2541
Local-gov        2093
?                1836
State-gov        1298
Self-emp-inc     1116
Federal-gov      960
Without-pay      14
Never-worked      7
Name: workclass, dtype: int64
HS-grad          10501
Some-college     7291
Bachelors        5355
Masters          1723
Assoc-voc        1382
11th             1175
Assoc-acdm       1067
10th             933
7th-8th          646
Prof-school      576
9th              514
12th             433
Doctorate        413
5th-6th          333
1st-4th          168
Preschool        51
Name: education, dtype: int64
```

Married-civ-spouse	14976
Never-married	10683
Divorced	4443
Separated	1025
Widowed	993
Married-spouse-absent	418
Married-AF-spouse	23
Name: marital_status, dtype: int64	
Prof-specialty	4140
Craft-repair	4099
Exec-managerial	4066
Adm-clerical	3770
Sales	3650
Other-service	3295
Machine-op-inspct	2002
?	1843
Transport-moving	1597
Handlers-cleaners	1370
Farming-fishing	994
Tech-support	928
Protective-serv	649
Priv-house-serv	149
Armed-Forces	9
Name: occupation, dtype: int64	
Husband	13193
Not-in-family	8305
Own-child	5068
Unmarried	3446
Wife	1568
Other-relative	981
Name: relationship, dtype: int64	
White	27816
Black	3124
Asian-Pac-Islander	1039
Amer-Indian-Eskimo	311
Other	271
Name: race, dtype: int64	
Male	21790
Female	10771
Name: sex, dtype: int64	
United-States	29170
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100



Cuba	95
England	90
Jamaica	81
South	80
China	75
Italy	73
Dominican-Republic	70
Vietnam	67
Guatemala	64
Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Trinidad&Tobago	19
Cambodia	19
Thailand	18
Laos	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native\_country, dtype: int64

<=50K      24720

>50K        7841

Name: income, dtype: int64

*# view frequency distribution of categorical variables*

for var in categorical:

    print(df[var].value\_counts()/np.float(len(df)))

Private	0.697030
Self-emp-not-inc	0.078038
Local-gov	0.064279
?	0.056386
State-gov	0.039864
Self-emp-inc	0.034274

```

Federal-gov      0.029483
Without-pay      0.000430
Never-worked     0.000215
Name: workclass, dtype: float64
HS-grad          0.322502
Some-college     0.223918
Bachelors        0.164461
Masters          0.052916
Assoc-voc        0.042443
11th             0.036086
Assoc-acdm       0.032769
10th             0.028654
7th-8th          0.019840
Prof-school      0.017690
9th              0.015786
12th             0.013298
Doctorate        0.012684
5th-6th          0.010227
1st-4th          0.005160
Preschool        0.001566
Name: education, dtype: float64
Married-civ-spouse 0.459937
Never-married      0.328092
Divorced           0.136452
Separated          0.031479
Widowed            0.030497
Married-spouse-absent 0.012837
Married-AF-spouse  0.000706
Name: marital_status, dtype: float64
Prof-specialty     0.127146
Craft-repair       0.125887
Exec-managerial    0.124873
Adm-clerical       0.115783
Sales              0.112097
Other-service      0.101195
Machine-op-inspct  0.061485
?                  0.056601
Transport-moving   0.049046
Handlers-cleaners  0.042075
Farming-fishing    0.030527
Tech-support       0.028500
Protective-serv    0.019932
Priv-house-serv    0.004576
Armed-Forces       0.000276
Name: occupation, dtype: float64
Husband            0.405178
Not-in-family      0.255060
Own-child          0.155646
Unmarried          0.105832

```

Wife	0.048156
Other-relative	0.030128
Name: relationship, dtype: float64	
White	0.854274
Black	0.095943
Asian-Pac-Islander	0.031909
Amer-Indian-Eskimo	0.009551
Other	0.008323
Name: race, dtype: float64	
Male	0.669205
Female	0.330795
Name: sex, dtype: float64	
United-States	0.895857
Mexico	0.019748
?	0.017905
Philippines	0.006081
Germany	0.004207
Canada	0.003716
Puerto-Rico	0.003501
El-Salvador	0.003255
India	0.003071
Cuba	0.002918
England	0.002764
Jamaica	0.002488
South	0.002457
China	0.002303
Italy	0.002242
Dominican-Republic	0.002150
Vietnam	0.002058
Guatemala	0.001966
Japan	0.001904
Poland	0.001843
Columbia	0.001812
Taiwan	0.001566
Haiti	0.001351
Iran	0.001321
Portugal	0.001136
Nicaragua	0.001044
Peru	0.000952
France	0.000891
Greece	0.000891
Ecuador	0.000860
Ireland	0.000737
Hong	0.000614
Trinidad&Tobago	0.000584
Cambodia	0.000584
Thailand	0.000553
Laos	0.000553
Yugoslavia	0.000491

```

Outlying-US(Guam-USVI-etc)    0.000430
Honduras                      0.000399
Hungary                      0.000399
Scotland                     0.000369
Holand-Netherlands           0.000031
Name: native_country, dtype: float64
<=50K      0.75919
>50K       0.24081
Name: income, dtype: float64

```

Now, we can see that there are several variables like `workclass`, `occupation` and `native_country` which contain missing values. Generally, the missing values are coded as `NaN` and python will detect them with the usual command of `df.isnull().sum()`.

But, in this case the missing values are coded as `?`. Python fail to detect these as missing values because it do not consider `?` as missing values. So, I have to replace `?` with `NaN` so that Python can detect these missing values.

I will explore these variables and replace `?` with `NaN`.

## Explore workclass variable

```

# check labels in workclass variable

df.workclass.unique()

array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
       'Local-gov', '?', 'Self-emp-inc', 'Without-pay', 'Never-
worked'],
      dtype=object)

# check frequency distribution of values in workclass variable

df.workclass.value_counts()

Private      22696
Self-emp-not-inc  2541
Local-gov    2093
?            1836
State-gov    1298
Self-emp-inc  1116
Federal-gov   960
Without-pay   14
Never-worked   7
Name: workclass, dtype: int64

```

We can see that there are 1836 values encoded as `?` in workclass variable. I will replace these `?` with `NaN`.

```
# replace '?' values in workclass variable with `NaN`

df['workclass'].replace('?', np.NaN, inplace=True)

# again check the frequency distribution of values in workclass variable

df.workclass.value_counts()

Private                22696
Self-emp-not-inc       2541
Local-gov              2093
State-gov              1298
Self-emp-inc           1116
Federal-gov            960
Without-pay            14
Never-worked           7
Name: workclass, dtype: int64
```

Now, we can see that there are no values encoded as ? in the `workclass` variable.

I will adopt similar approach with `occupation` and `native_country` column.

## Explore occupation variable

```
# check labels in occupation variable

df.occupation.unique()

array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
       'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
       'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
       'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
       'Priv-house-serv'], dtype=object)

# check frequency distribution of values in occupation variable

df.occupation.value_counts()

Prof-specialty         4140
Craft-repair           4099
Exec-managerial        4066
Adm-clerical           3770
Sales                  3650
Other-service          3295
Machine-op-inspct      2002
?                      1843
Transport-moving       1597
Handlers-cleaners      1370
Farming-fishing        994
```

```
Tech-support          928
Protective-serv       649
Priv-house-serv       149
Armed-Forces          9
Name: occupation, dtype: int64
```

We can see that there are 1843 values encoded as `?` in `occupation` variable. I will replace these `?` with `NaN`.

```
# replace '?' values in occupation variable with `NaN`
df['occupation'].replace('?', np.NaN, inplace=True)

# again check the frequency distribution of values in occupation
variable

df.occupation.value_counts()

Prof-specialty        4140
Craft-repair          4099
Exec-managerial       4066
Adm-clerical          3770
Sales                 3650
Other-service         3295
Machine-op-inspct     2002
Transport-moving      1597
Handlers-cleaners     1370
Farming-fishing       994
Tech-support          928
Protective-serv       649
Priv-house-serv       149
Armed-Forces          9
Name: occupation, dtype: int64
```

## Explore native\_country variable

```
# check labels in native_country variable

df.native_country.unique()

array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
      'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada',
      'Germany',
      'Iran', 'Philippines', 'Italy', 'Poland', 'Columbia',
      'Cambodia',
      'Thailand', 'Ecuador', 'Laos', 'Taiwan', 'Haiti', 'Portugal',
      'Dominican-Republic', 'El-Salvador', 'France', 'Guatemala',
      'China', 'Japan', 'Yugoslavia', 'Peru',
      'Outlying-US(Guam-USVI-etc)', 'Scotland', 'Trinidad&Tobago',
```

```
'Greece', 'Nicaragua', 'Vietnam', 'Hong', 'Ireland', 'Hungary',  
'Holand-Netherlands'], dtype=object)
```

```
# check frequency distribution of values in native_country variable
```

```
df.native_country.value_counts()
```

United-States	29170
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100
Cuba	95
England	90
Jamaica	81
South	80
China	75
Italy	73
Dominican-Republic	70
Vietnam	67
Guatemala	64
Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Trinidad&Tobago	19
Cambodia	19
Thailand	18
Laos	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native\_country, dtype: int64

We can see that there are 583 values encoded as ? in native\_country variable. I will replace these ? with NaN.

```
# replace '?' values in native_country variable with `NaN`
df['native_country'].replace('?', np.NaN, inplace=True)

# again check the frequency distribution of values in native_country variable
df.native_country.value_counts()
```

United-States	29170
Mexico	643
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100
Cuba	95
England	90
Jamaica	81
South	80
China	75
Italy	73
Dominican-Republic	70
Vietnam	67
Guatemala	64
Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Trinidad&Tobago	19
Cambodia	19
Thailand	18
Laos	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13



```
Hungary          13
Scotland         12
Holand-Netherlands 1
Name: native_country, dtype: int64
```

Check missing values in categorical variables again

```
df[categorical].isnull().sum()
```

```
workclass      1836
education       0
marital_status  0
occupation     1843
relationship    0
race            0
sex            0
native_country  583
income         0
dtype: int64
```

Now, we can see that `workclass`, `occupation` and `native_country` variable contains missing values.

## Number of labels: cardinality

The number of labels within a categorical variable is known as **cardinality**. A high number of labels within a variable is known as **high cardinality**. High cardinality may pose some serious problems in the machine learning model. So, I will check for high cardinality.

```
# check for cardinality in categorical variables
```

```
for var in categorical:
```

```
    print(var, ' contains ', len(df[var].unique()), ' labels')
```

```
workclass contains 9 labels
education contains 16 labels
marital_status contains 7 labels
occupation contains 15 labels
relationship contains 6 labels
race contains 5 labels
sex contains 2 labels
native_country contains 42 labels
income contains 2 labels
```

We can see that `native_country` column contains relatively large number of labels as compared to other columns. I will check for cardinality after train-test split.

## Explore Numerical Variables

```
# find numerical variables

numerical = [var for var in df.columns if df[var].dtype!='O']

print('There are {} numerical variables\n'.format(len(numerical)))

print('The numerical variables are :', numerical)

There are 6 numerical variables

The numerical variables are : ['age', 'fnlwgt', 'education_num',
'capital_gain', 'capital_loss', 'hours_per_week']

# view the numerical variables

df[numerical].head()
```

	age	fnlwgt	education_num	capital_gain	capital_loss
hours_per_week					
0	39	77516	13	2174	0
40					
1	50	83311	13	0	0
13					
2	38	215646	9	0	0
40					
3	53	234721	7	0	0
40					
4	28	338409	13	0	0
40					

## Summary of numerical variables

- There are 6 numerical variables.
- These are given by age, fnlwgt, education\_num, capital\_gain, capital\_loss and hours\_per\_week.
- All of the numerical variables are of discrete data type.

## Explore problems within numerical variables

Now, I will explore the numerical variables.

## Missing values in numerical variables

```
# check missing values in numerical variables

df[numerical].isnull().sum()
```

```
age          0
fnlwgt       0
education_num 0
capital_gain 0
capital_loss 0
hours_per_week 0
dtype: int64
```

We can see that all the 6 numerical variables do not contain missing values.

## 8. Declare feature vector and target variable

```
X = df.drop(['income'], axis=1)
y = df['income']
```

## 9. Split data into separate training and test set

```
# split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)

# check the shape of X_train and X_test

X_train.shape, X_test.shape

((22792, 14), (9769, 14))
```

## 10. Feature Engineering

**Feature Engineering** is the process of transforming raw data into useful features that help us to understand our model better and increase its predictive power. I will carry out feature engineering on different types of variables.

First, I will display the categorical and numerical variables again separately.

```
# check data types in X_train

X_train.dtypes

age          int64
workclass    object
fnlwgt       int64
education    object
education_num int64
marital_status object
```

```

occupation      object
relationship     object
race            object
sex            object
capital_gain    int64
capital_loss    int64
hours_per_week  int64
native_country  object
dtype: object

# display categorical variables

categorical = [col for col in X_train.columns if X_train[col].dtypes
== 'O']

categorical

['workclass',
 'education',
 'marital_status',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native_country']

# display numerical variables

numerical = [col for col in X_train.columns if X_train[col].dtypes !=
'O']

numerical

['age',
 'fnlwgt',
 'education_num',
 'capital_gain',
 'capital_loss',
 'hours_per_week']

```

## Engineering missing values in categorical variables

```

# print percentage of missing values in the categorical variables in
training set

X_train[categorical].isnull().mean()

workclass      0.055985
education      0.000000
marital_status 0.000000
occupation     0.056072

```

```

relationship      0.000000
race              0.000000
sex              0.000000
native_country    0.018164
dtype: float64

# print categorical variables with missing data

for col in categorical:
    if X_train[col].isnull().mean()>0:
        print(col, (X_train[col].isnull().mean()))

workclass 0.055984555984555984
occupation 0.05607230607230607
native_country 0.018164268164268166

# impute missing categorical variables with most frequent value

for df2 in [X_train, X_test]:
    df2['workclass'].fillna(X_train['workclass'].mode()[0],
inplace=True)
    df2['occupation'].fillna(X_train['occupation'].mode()[0],
inplace=True)
    df2['native_country'].fillna(X_train['native_country'].mode()[0],
inplace=True)

# check missing values in categorical variables in X_train

X_train[categorical].isnull().sum()

workclass      0
education      0
marital_status 0
occupation     0
relationship   0
race           0
sex            0
native_country 0
dtype: int64

# check missing values in categorical variables in X_test

X_test[categorical].isnull().sum()

workclass      0
education      0
marital_status 0
occupation     0
relationship   0
race           0
sex            0

```

```
native_country    0
dtype: int64
```

As a final check, I will check for missing values in X\_train and X\_test.

```
# check missing values in X_train
```

```
X_train.isnull().sum()
```

```
age            0
workclass      0
fnlwgt        0
education      0
education_num  0
marital_status 0
occupation     0
relationship   0
race           0
sex            0
capital_gain   0
capital_loss   0
hours_per_week 0
native_country 0
dtype: int64
```

```
# check missing values in X_test
```

```
X_test.isnull().sum()
```

```
age            0
workclass      0
fnlwgt        0
education      0
education_num  0
marital_status 0
occupation     0
relationship   0
race           0
sex            0
capital_gain   0
capital_loss   0
hours_per_week 0
native_country 0
dtype: int64
```

We can see that there are no missing values in X\_train and X\_test.

## Encode categorical variables

```
# print categorical variables
```

```
categorical
```

```
['workclass',  
'education',  
'marital_status',  
'occupation',  
'relationship',  
'race',  
'sex',  
'native_country']
```

```
X_train[categorical].head()
```

	workclass	education	marital_status	occupation \
32098	Private	HS-grad	Married-civ-spouse	Craft-repair
25206	State-gov	HS-grad	Divorced	Adm-clerical
23491	Private	Some-college	Married-civ-spouse	Sales
12367	Private	HS-grad	Never-married	Craft-repair
7054	Private	7th-8th	Never-married	Craft-repair

	relationship	race	sex	native_country
32098	Husband	White	Male	United-States
25206	Unmarried	White	Female	United-States
23491	Husband	White	Male	United-States
12367	Not-in-family	White	Male	Guatemala
7054	Not-in-family	White	Male	Germany

```
# import category encoders
```

```
import category_encoders as ce
```

```
# encode remaining variables with one-hot encoding
```

```
encoder = ce.OneHotEncoder(cols=['workclass', 'education',  
'marital_status', 'occupation', 'relationship',  
'race', 'sex', 'native_country'])
```

```
X_train = encoder.fit_transform(X_train)
```

```
X_test = encoder.transform(X_test)
```

```
X_train.head()
```

	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5
32098	1	0	0	0	0
25206	0	1	0	0	0

23491	1	0	0	0	0
12367	1	0	0	0	0
7054	1	0	0	0	0

	workclass_6	workclass_7	workclass_8	workclass_-1
education_1 \				
32098	0	0	0	0
1				
25206	0	0	0	0
1				
23491	0	0	0	0
0				
12367	0	0	0	0
1				
7054	0	0	0	0
0				

	...	native_country_39	native_country_40	\
32098	...	0	0	
25206	...	0	0	
23491	...	0	0	
12367	...	0	0	
7054	...	0	0	

	native_country_41	native_country_-1	age	fnlwgt
education_num \				
32098	0	0	45	170871
9				
25206	0	0	47	108890
9				
23491	0	0	48	187505
10				
12367	0	0	29	145592
9				
7054	0	0	23	203003
4				

	capital_gain	capital_loss	hours_per_week
32098	7298	0	60
25206	1831	0	38
23491	0	0	50
12367	0	0	40
7054	0	0	25

[5 rows x 113 columns]



```
X_train.shape  
(22792, 113)
```

We can see that from the initial 14 columns, we now have 113 columns.

Similarly, I will take a look at the `X_test` set.

```
X_test.head()
```

	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5
22278	1	0	0	0	0
8950	1	0	0	0	0
7838	1	0	0	0	0
16505	1	0	0	0	0
19140	1	0	0	0	0

	workclass_6	workclass_7	workclass_8	workclass_-1
22278	0	0	0	0
8950	0	0	0	0
7838	0	0	0	0
16505	0	0	0	0
19140	0	0	0	0

	...	native_country_39	native_country_40	\
22278	...	0	0	
8950	...	0	0	
7838	...	0	0	
16505	...	0	0	
19140	...	0	0	

	native_country_41	native_country_-1	age	fnlwgt
22278	0	0	27	177119
8950	0	0	27	216481
7838	0	0	25	256263

```

12
16505          0          0  46  147640
3
19140          0          0  45  172822
7

```

```

      capital_gain  capital_loss  hours_per_week
22278           0           0           44
8950            0           0           40
7838            0           0           40
16505            0          1902           40
19140            0          2824           76

```

```
[5 rows x 113 columns]
```

```
X_test.shape
```

```
(9769, 113)
```

We now have training and testing set ready for model building. Before that, we should map all the feature variables onto the same scale. It is called **feature scaling**. I will do it as follows.

## 11. Feature Scaling

```
cols = X_train.columns
```

```
from sklearn.preprocessing import RobustScaler
```

```
scaler = RobustScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
X_train = pd.DataFrame(X_train, columns=[cols])
```

```
X_test = pd.DataFrame(X_test, columns=[cols])
```

```
X_train.head()
```

```

      workclass_1 workclass_2 workclass_3 workclass_4 workclass_5
workclass_6 \
0           0.0           0.0           0.0           0.0           0.0
0.0
1          -1.0           1.0           0.0           0.0           0.0
0.0
2           0.0           0.0           0.0           0.0           0.0
0.0
3           0.0           0.0           0.0           0.0           0.0
0.0
4           0.0           0.0           0.0           0.0           0.0

```

```

0.0
workclass_7 workclass_8 workclass_-1 education_1 ... \
0 0.0 0.0 0.0 1.0 ...
1 0.0 0.0 0.0 1.0 ...
2 0.0 0.0 0.0 0.0 ...
3 0.0 0.0 0.0 1.0 ...
4 0.0 0.0 0.0 0.0 ...

native_country_39 native_country_40 native_country_41
native_country_-1 \
0 0.0 0.0 0.0
0.0
1 0.0 0.0 0.0
0.0
2 0.0 0.0 0.0
0.0
3 0.0 0.0 0.0
0.0
4 0.0 0.0 0.0
0.0

age fnlwgt education_num capital_gain capital_loss
hours_per_week
0 0.40 -0.058906 -0.333333 7298.0 0.0
4.0
1 0.50 -0.578076 -0.333333 1831.0 0.0 -
0.4
2 0.55 0.080425 0.000000 0.0 0.0
2.0
3 -0.40 -0.270650 -0.333333 0.0 0.0
0.0
4 -0.70 0.210240 -2.000000 0.0 0.0 -
3.0

[5 rows x 113 columns]

```

We now have `X_train` dataset ready to be fed into the Gaussian Naive Bayes classifier. I will do it as follows.

## 12. Model training

```

# train a Gaussian Naive Bayes classifier on the training set
from sklearn.naive_bayes import GaussianNB

# instantiate the model
gnb = GaussianNB()

```

```
# fit the model
gnb.fit(X_train, y_train)

GaussianNB(priors=None, var_smoothing=1e-09)
```

## 13. Predict the test set results

```
y_pred = gnb.predict(X_test)

y_pred
array(['<=50K', '<=50K', '>50K', ..., '>50K', '<=50K', '<=50K'],
      dtype='<U5')
```

## 14. Check accuracy score

```
from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test,
y_pred)))

Model accuracy score: 0.8083
```

Here, **y\_test** are the true class labels and **y\_pred** are the predicted class labels in the test-set.

### Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
y_pred_train = gnb.predict(X_train)

y_pred_train
array(['>50K', '<=50K', '>50K', ..., '<=50K', '>50K', '<=50K'],
      dtype='<U5')

print('Training-set accuracy score: {0:0.4f}'.
      format(accuracy_score(y_train, y_pred_train)))

Training-set accuracy score: 0.8067
```

### Check for overfitting and underfitting

```
# print the scores on training and test set

print('Training set score: {:.4f}'.format(gnb.score(X_train,
y_train)))

print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))
```

```
Training set score: 0.8067
Test set score: 0.8083
```

The training-set accuracy score is 0.8067 while the test-set accuracy to be 0.8083. These two values are quite comparable. So, there is no sign of overfitting.

## Compare model accuracy with null accuracy

So, the model accuracy is 0.8083. But, we cannot say that our model is very good based on the above accuracy. We must compare it with the **null accuracy**. Null accuracy is the accuracy that could be achieved by always predicting the most frequent class.

So, we should first check the class distribution in the test set.

```
# check class distribution in test set
```

```
y_test.value_counts()
```

```
<=50K    7407
>50K     2362
Name: income, dtype: int64
```

We can see that the occurrences of most frequent class is 7407. So, we can calculate null accuracy by dividing 7407 by total number of occurrences.

```
# check null accuracy score
```

```
null_accuracy = (7407/(7407+2362))
```

```
print('Null accuracy score: {0:0.4f}'.format(null_accuracy))
```

```
Null accuracy score: 0.7582
```

We can see that our model accuracy score is 0.8083 but null accuracy score is 0.7582. So, we can conclude that our Gaussian Naive Bayes Classification model is doing a very good job in predicting the class labels.

Now, based on the above analysis we can conclude that our classification model accuracy is very good. Our model is doing a very good job in terms of predicting the class labels.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called **Confusion matrix** that comes to our rescue.

## 15. Confusion matrix

A confusion matrix is a tool for summarizing the performance of a classification algorithm. A confusion matrix will give us a clear picture of classification model performance and the types of

errors produced by the model. It gives us a summary of correct and incorrect predictions broken down by each category. The summary is represented in a tabular form.

Four types of outcomes are possible while evaluating a classification model performance. These four outcomes are described below:-

**True Positives (TP)** – True Positives occur when we predict an observation belongs to a certain class and the observation actually belongs to that class.

**True Negatives (TN)** – True Negatives occur when we predict an observation does not belong to a certain class and the observation actually does not belong to that class.

**False Positives (FP)** – False Positives occur when we predict an observation belongs to a certain class but the observation actually does not belong to that class. This type of error is called **Type I error**.

**False Negatives (FN)** – False Negatives occur when we predict an observation does not belong to a certain class but the observation actually belongs to that class. This is a very serious error and it is called **Type II error**.

These four outcomes are summarized in a confusion matrix given below.

```
# Print the Confusion Matrix and slice it into four pieces
```

```
from sklearn.metrics import confusion_matrix
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
print('Confusion matrix\n\n', cm)
```

```
print('\nTrue Positives(TP) = ', cm[0,0])
```

```
print('\nTrue Negatives(TN) = ', cm[1,1])
```

```
print('\nFalse Positives(FP) = ', cm[0,1])
```

```
print('\nFalse Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix
```

```
[[5999 1408]
 [ 465 1897]]
```

```
True Positives(TP) = 5999
```

```
True Negatives(TN) = 1897
```

```
False Positives(FP) = 1408
```


```
False Negatives(FN) = 465
```

The confusion matrix shows  $5999 + 1897 = 7896$  correct predictions and  $1408 + 465 = 1873$  incorrect predictions.

In this case, we have

- True Positives (Actual Positive:1 and Predict Positive:1) - 5999
- True Negatives (Actual Negative:0 and Predict Negative:0) - 1897
- False Positives (Actual Negative:0 but Predict Positive:1) - 1408 (Type I error)
- False Negatives (Actual Positive:1 but Predict Negative:0) - 465 (Type II error)

```
# visualize confusion matrix with seaborn heatmap  
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1',  
                                           'Actual Negative:0'],  
                          index=['Predict Positive:1', 'Predict  
Negative:0'])  
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')  
<matplotlib.axes._subplots.AxesSubplot at 0xf202a9c6a0>
```



## 16. Classification metrices

### Classification Report

**Classification report** is another way to evaluate the classification model performance. It displays the **precision**, **recall**, **f1** and **support** scores for the model. I have described these terms in later.

We can print a classification report as follows:-

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.93	0.81	0.86	7407
>50K	0.57	0.80	0.67	2362
micro avg	0.81	0.81	0.81	9769
macro avg	0.75	0.81	0.77	9769
weighted avg	0.84	0.81	0.82	9769

### Classification accuracy

```
TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]

# print classification accuracy

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy :
{0:0.4f}'.format(classification_accuracy))

Classification accuracy : 0.8083
```

### Classification error

```
# print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))

Classification error : 0.1917
```



## Precision

**Precision** can be defined as the percentage of correctly predicted positive outcomes out of all the predicted positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true and false positives (TP + FP).

So, **Precision** identifies the proportion of correctly predicted positive outcome. It is more concerned with the positive class than the negative class.

Mathematically, precision can be defined as the ratio of TP to (TP + FP).

```
# print precision score
precision = TP / float(TP + FP)

print('Precision : {0:0.4f}'.format(precision))
Precision : 0.8099
```

## Recall

Recall can be defined as the percentage of correctly predicted positive outcomes out of all the actual positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true positives and false negatives (TP + FN). **Recall** is also called **Sensitivity**.

**Recall** identifies the proportion of correctly predicted actual positives.

Mathematically, recall can be given as the ratio of TP to (TP + FN).

```
recall = TP / float(TP + FN)

print('Recall or Sensitivity : {0:0.4f}'.format(recall))
Recall or Sensitivity : 0.9281
```

## True Positive Rate

**True Positive Rate** is synonymous with **Recall**.

```
true_positive_rate = TP / float(TP + FN)

print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
True Positive Rate : 0.9281
```

## False Positive Rate

```
false_positive_rate = FP / float(FP + TN)
```

```
print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))
False Positive Rate : 0.4260
```

## Specificity

```
specificity = TN / (TN + FP)
print('Specificity : {0:0.4f}'.format(specificity))
Specificity : 0.5740
```

## f1-score

**f1-score** is the weighted harmonic mean of precision and recall. The best possible **f1-score** would be 1.0 and the worst would be 0.0. **f1-score** is the harmonic mean of precision and recall. So, **f1-score** is always lower than accuracy measures as they embed precision and recall into their computation. The weighted average of **f1-score** should be used to compare classifier models, not global accuracy.

## Support

**Support** is the actual number of occurrences of the class in our dataset.

## 17. Calculate class probabilities

```
# print the first 10 predicted probabilities of two classes- 0 and 1
y_pred_prob = gnb.predict_proba(X_test)[0:10]
y_pred_prob
array([[9.99999426e-01, 5.74152436e-07],
       [9.99687907e-01, 3.12093456e-04],
       [1.54405602e-01, 8.45594398e-01],
       [1.73624321e-04, 9.99826376e-01],
       [8.20121011e-09, 9.99999992e-01],
       [8.76844580e-01, 1.23155420e-01],
       [9.99999927e-01, 7.32876705e-08],
       [9.99993460e-01, 6.53998797e-06],
       [9.87738143e-01, 1.22618575e-02],
       [9.99999996e-01, 4.01886317e-09]])
```

## Observations

- In each row, the numbers sum to 1.
- There are 2 columns which correspond to 2 classes - `<=50K` and `>50K`.
  - Class 0 => `<=50K` - Class that a person makes less than equal to 50K.

- Class 1 => >50K - Class that a person makes more than 50K.
- Importance of predicted probabilities
  - We can rank the observations by probability of whether a person makes less than or equal to 50K or more than 50K.
- predict\_proba process
  - Predicts the probabilities
  - Choose the class with the highest probability
- Classification threshold level
  - There is a classification threshold level of 0.5.
  - Class 0 => <=50K - probability of salary less than or equal to 50K is predicted if probability < 0.5.
  - Class 1 => >50K - probability of salary more than 50K is predicted if probability > 0.5.

```
# store the probabilities in dataframe
```

```
y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of - <=50K', 'Prob of - >50K'])
```

```
y_pred_prob_df
```

	Prob of - <=50K	Prob of - >50K
0	9.999994e-01	5.741524e-07
1	9.996879e-01	3.120935e-04
2	1.544056e-01	8.455944e-01
3	1.736243e-04	9.998264e-01
4	8.201210e-09	1.000000e+00
5	8.768446e-01	1.231554e-01
6	9.999999e-01	7.328767e-08
7	9.999935e-01	6.539988e-06
8	9.877381e-01	1.226186e-02
9	1.000000e+00	4.018863e-09

```
# print the first 10 predicted probabilities for class 1 - Probability of >50K
```

```
gnb.predict_proba(X_test)[0:10, 1]
```

```
array([5.74152436e-07, 3.12093456e-04, 8.45594398e-01, 9.99826376e-01,
        9.9999992e-01, 1.23155420e-01, 7.32876705e-08, 6.53998797e-06,
        1.22618575e-02, 4.01886317e-09])
```

```
# store the predicted probabilities for class 1 - Probability of >50K
y_pred1 = gnb.predict_proba(X_test)[: , 1]
# plot histogram of predicted probabilities


# adjust the font size
plt.rcParams['font.size'] = 12

# plot histogram with 10 bins
plt.hist(y_pred1, bins = 10)

# set the title of predicted probabilities
plt.title('Histogram of predicted probabilities of salaries >50K')

# set the x-axis limit
plt.xlim(0,1)

# set the title
plt.xlabel('Predicted probabilities of salaries >50K')
plt.ylabel('Frequency')
Text(0,0.5,'Frequency')
```



## Observations

- We can see that the above histogram is highly positive skewed.
- The first column tell us that there are approximately 5700 observations with probability between 0.0 and 0.1 whose salary is  $\leq 50K$ .
- There are relatively small number of observations with probability  $> 0.5$ .
- So, these small number of observations predict that the salaries will be  $> 50K$ .
- Majority of observations predcit that the salaries will be  $\leq 50K$ .

## 18. ROC - AUC

### ROC Curve

Another tool to measure the classification model performance visually is **ROC Curve**. ROC Curve stands for **Receiver Operating Characteristic Curve**. An **ROC Curve** is a plot which shows the performance of a classification model at various classification threshold levels.

The **ROC Curve** plots the **True Positive Rate (TPR)** against the **False Positive Rate (FPR)** at various threshold levels.

**True Positive Rate (TPR)** is also called **Recall**. It is defined as the ratio of TP to (TP + FN).

**False Positive Rate (FPR)** is defined as the ratio of FP to (FP + TN).

In the ROC Curve, we will focus on the TPR (True Positive Rate) and FPR (False Positive Rate) of a single point. This will give us the general performance of the ROC curve which consists of the TPR and FPR at various threshold levels. So, an ROC Curve plots TPR vs FPR at different classification threshold levels. If we lower the threshold levels, it may result in more items being classified as positive. It will increase both True Positives (TP) and False Positives (FP).

```
# plot ROC Curve

from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = '>50K')

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)

plt.plot([0,1], [0,1], 'k--' )

plt.rcParams['font.size'] = 12


plt.title('ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries')

plt.xlabel('False Positive Rate (1 - Specificity)')

plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```

ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries



ROC curve help us to choose a threshold level that balances sensitivity and specificity for a particular context.

## ROC AUC

**ROC AUC** stands for **Receiver Operating Characteristic - Area Under Curve**. It is a technique to compare classifier performance. In this technique, we measure the area under the curve (AUC). A perfect classifier will have a ROC AUC equal to 1, whereas a purely random classifier will have a ROC AUC equal to 0.5.

So, **ROC AUC** is the percentage of the ROC plot that is underneath the curve.

```
# compute ROC AUC

from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC : {:.4f}'.format(ROC_AUC))

ROC AUC : 0.8941
```

## Interpretation

- ROC AUC is a single number summary of classifier performance. The higher the value, the better the classifier.
- ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a good job in predicting whether it will rain tomorrow or not.

```
# calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(gnb, X_train, y_train, cv=5,
scoring='roc_auc').mean()

print('Cross validated ROC AUC :
{:.4f}'.format(Cross_validated_ROC_AUC))

Cross validated ROC AUC : 0.8938
```

## 19. k-Fold Cross Validation

```
# Applying 10-Fold Cross Validation

from sklearn.model_selection import cross_val_score

scores = cross_val_score(gnb, X_train, y_train, cv = 10,
scoring='accuracy')
```

```
print('Cross-validation scores:{}'.format(scores))
```

```
Cross-validation scores:[0.81359649 0.80438596 0.81184211 0.80517771  
0.79640193 0.79684072  
0.81044318 0.81175954 0.80210619 0.81035996]
```

We can summarize the cross-validation accuracy by calculating its mean.

```
# compute Average cross-validation score
```

```
print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

```
Average cross-validation score: 0.8063
```

## Interpretation

- Using the mean cross-validation, we can conclude that we expect the model to be around 80.63% accurate on average.
- If we look at all the 10 scores produced by the 10-fold cross-validation, we can also conclude that there is a relatively small variance in the accuracy between folds, ranging from 81.35% accuracy to 79.64% accuracy. So, we can conclude that the model is independent of the particular folds used for training.
- Our original model accuracy is 0.8083, but the mean cross-validation accuracy is 0.8063. So, the 10-fold cross-validation accuracy does not result in performance improvement for this model.

## 20. Results and conclusion

1. In this project, I build a Gaussian Naïve Bayes Classifier model to predict whether a person makes over 50K a year. The model yields a very good performance as indicated by the model accuracy which was found to be 0.8083.
2. The training-set accuracy score is 0.8067 while the test-set accuracy to be 0.8083. These two values are quite comparable. So, there is no sign of overfitting.
3. I have compared the model accuracy score which is 0.8083 with null accuracy score which is 0.7582. So, we can conclude that our Gaussian Naïve Bayes classifier model is doing a very good job in predicting the class labels.
4. ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a very good job in predicting whether a person makes over 50K a year.
5. Using the mean cross-validation, we can conclude that we expect the model to be around 80.63% accurate on average.
6. If we look at all the 10 scores produced by the 10-fold cross-validation, we can also conclude that there is a relatively small variance in the accuracy between folds, ranging from 81.35% accuracy to 79.64% accuracy. So, we can conclude that the model is independent of the particular folds used for training.



7. Our original model accuracy is 0.8083, but the mean cross-validation accuracy is 0.8063. So, the 10-fold cross-validation accuracy does not result in performance improvement for this model.